

# COMP3031: Parallel and Concurrent Programming

Qiong Luo

HKUST

Concurrency can happen at different levels:

- Instruction level - executing two or more machine instructions simultaneously.
- Statement level.
- Unit level - executing two or more subprograms simultaneously.
- Program level - executing two or more programs simultaneously.

and can occur either

- physically on separate processors or
- logically through a single processor.

# Motivations for Studying Concurrency

- Concurrency is a natural way to conceptualize and solve many problems
  - ▶ Concurrency is ubiquitous, and many problems lend themselves naturally to concurrency in much the same way that recursion is a natural way to solve a problem.
  - ▶ This is particularly true in simulation systems: the systems being simulated contain many independent entities.
  - ▶ Think of cars in a traffic simulation system, airplanes in a air-control system,...
- Computing devices with multiple processors are becoming more common, and we need software to make effective use of these devices.

# Subprogram-Level Concurrency

## Basic Concepts:

- Task or process: a piece of code that can be run concurrently.
  - ▶ Heavyweight tasks: have their own address space and run-time stacks.
  - ▶ Lightweight tasks: have to share the same address space on the machine/processor where they are run.
  - ▶ States of a process: new; runnable or ready; running; blocked; dead.
- Synchronization: a mechanism that controls the order in which tasks execute
  - ▶ Cooperation synchronization: task A's job depends on task B.
  - ▶ Competition synchronization: task A and B compete for a same resource.
- Scheduler: decides which task to run (where there are more ready tasks than processors).

# Examples

- A global variable  $x$ , and two tasks - A:  $x = x++$ ; B:  $x = 2x$ . How many possible results if they are performed concurrently:
- Producer-consumer problem: a number of producers produce some products and put them into a queue, and a number of consumers take them to consume from the queue. How should they be synchronized.

# Semaphores

Semaphores are special data structures devised by Dijkstra in 1965 for synchronizing concurrent tasks. In modern terms, one could think of them as a class with the following components:

- A private data member counter of int type, and a queue of tasks;
- two public methods:
  - ▶ wait - if `counter > 0` then `counter--` else put the caller in the queue.
  - ▶ release - if the queue is empty then `counter++` else dequeue a waiting task.

# Semaphores: Cooperation Synchronization

Allow simultaneous access to the VALUE queue:

```
semaphore fullspots, emptyspots;
fullspots.counter = 0;
emptyspots.counter = BUFLen;
task producer;
    while (true) {
        -- produce VALUE --
        wait(emptyspots);
        DEPOSIT(VALUE);
        release(fullspots); }
task consumer;
    while (true) {
        wait(fullspots);
        FETCH(VALUE);
        release(emptyspots);
        -- consume VALUE -- }
```

# Semaphores: Competition Synchronizations

Only one task can access VALUE queue at any time:

```
semaphore fullspots, emptyspots, access;  
fullspots.counter = 0;  
emptyspots.counter = BUFLen;  
access.counter = 1;  
task producer;  
    while (true) {  
        -- produce VALUE --  
        wait(emptyspots);  
        wait(access); // wait for access to the queue  
        DEPOSIT(VALUE);  
        release(access); //relinquish access  
        release(fullspots); }
```



```
task consumer;
  while (true) {
    wait(fullspots);
    wait(access);
    FETCH(VALUE);
    release(access);
    release(emptyspots);
    -- consume VALUE -- }
```

# Java Threads

- In Java, tasks are called threads. They are lightweight tasks, meaning that they can only be run on the same machine.
- Java has a build-in class called Thread, from which user-defined threads are derived.
- Threads are initiated by "start()" method, which in turn calls the "run()" method.
- start() returns immediately, but run() is the life-span of the thread.
- Some other methods in Thread:
  - ▶ yield() - voluntarily surrender the execution of the thread;
  - ▶ sleep(n) - suspend the thread for n milliseconds;
  - ▶ thr.join(n) - wait for n millisecond for the thread thr to complete;

# Competition Synchronization

In Java, competition synchronization is done by "synchronized methods" or "synchronized statements": in the following, "synchronized" is a reserved word

- "synchronized method()" locks the object until the method is finished;
- during the execution of method(), no other synchronized method is allowed to access the object.

```
class ManageBuf {  
    ...  
    public synchronized void deposit(int item) { ... }  
    public synchronized int fetch() { ... }  
    ... }
```

- "synchronized(expression) statement" locks the object specified by "expression" when "statement" is executed.

# Cooperation Synchronization

- Cooperation synchronization in Java is done with "wait", "notify", and "notifyAll" methods in Object - the root class of all Java classes.
- Every object has a wait list of all the threads that have called "wait" on that object.
- The "notify" method is called to free one of the waiting threads, randomly.
- The "notifyAll" method awakens all the threads on the object's waiting list.
- The "wait" method can throw "InterruptedException", so any code that uses "wait" must also catch this exception:

```
try {  
    while (!condition)  
        wait();  
    -- codes after 'condition' becomes true  
} catch (InterruptedException myProblem) { ... }
```

## Producer-Consumer in Java

```
//A circular queue for storing int values with synchronized
//methods for inserting and removing values from the queue
class Queue {
    private int [] que;
    private int nextIn, nextOut, filled, queSize;

    public Queue(int size) { que = new int [size];
        filled = 0; nextIn = 0; nextOut = 0; queSize = size; }

    public synchronized void deposit (int item) {
        try {
            while (filled == queSize) wait();
            que [nextIn] = item; nextIn = (nextIn++ % queSize);
            filled++; notifyAll(); }
        catch (InterruptedException e) {} }
```

## Producer-Consumer in Java (Cont'd)

```
public synchronized int fetch ( ) {
    int item = 0;
    try { while (filled == 0) wait();
        item=que[nextOut]; nextOut=(nextOut++ % queSize);
        filled--; notifyAll(); }
    catch (InterruptedException e) {}
    return item; }
} // end of Queue class

class Producer extends Thread {
    private Queue buffer;
    public Producer(Queue que) { buffer = que; }
    public void run() {
        int new_item;
        while (true) { //-- create a new_item
            buffer.deposit(new_item); }}
```

## Producer-Consumer in Java (Cont'd)

```
class Consumer extends Thread {  
    private Queue buffer;  
    public Consumer(Queue que) { buffer = que; }  
    public void run() {  
        int stored_item;  
        while (true) {  
            stored_item = buffer.fetch();  
            //-- consume stored_item  
        }  
    }  
}
```

```
Queue buff1 = new Queue(100);  
Producer prod1 = new Producer(buff1);  
Consumer cons1 = new Consumer(buff1);  
prod1.start(); cons1.start();
```

# Statement-Level Concurrency

- Statement-level concurrency is used to inform the compiler of ways it can map the program onto a multiprocessor machine.
- CUDA = a programming language for GPU - later.
- HPF - high-performance Fortran (1993), an extension of Fortran 90 to include statements on how concurrent programs should be run on a multiprocessor machine.
- These include specifications about the number of processors, statements about how data should be distributed over the memories of these processors, and the alignments of data.



# HPF Statements

- `!HPF$ PROCESSOR procs (n)` specifies there are `n` processors available.
- `!HPF$ DISTRIBUTE (kind) ONTO procs :: lists:`
  - ▶ Here 'kind' can be either `BLOCK` or `CYCLIC`, and 'lists' is the names of the array variables that are to be distributed.
  - ▶ `BLOCK` means that each array is to be divided into `n` (the number of processors) groups consisting of contiguous collections of array elements, and distributed over the memories of the `n` processors.  
Example: a 500 element array `BLOCK` distribute over 5 processors means that the first 100 elements on the first processor, the second 100 elements on the second processor, etc..
  - ▶ `CYCLIC` means to distribute the 1st element on the 1st processor, 2nd element on the 2nd processor, ..., `nth` element on the `nth` processor, `(n+1)th` element on the 1st processor, ...
- `!HPF$ ALIGN list1(index) WITH list2(index+1)` says that for all values of `index`, store `list1(index)` and `list2(index+1)` in the memory of the same processor.

## Example

```
REAL list1 (1000), list2(1000)
INTEGER list3(500), list4(501)
!HPF$ PROCESSOR procs (10)
!HPF$ DISTRIBUTE (BLOCK) ONTO procs :: list1, list2
!HPF$ ALIGN list3(index) WITH list4(index+1)
FORALL (index = 1:1000) list1(index) = list2(index)+1.0
FORALL (index = 1:500) list3(index) = list4(index+1)
```

FORALL statement specifies a collection of statements to be executed concurrently.