

COMP 3311 Database Management Systems Spring 2015

Lab 6. Oracle indexing, Oracle clustering and more about PL/SQL.

Objectives of the Lab

- After this lab you should be able to
 - know how to create an index for a table in Oracle,
 - know how to create a cluster for tables in Oracle,
 - know more about PL/SQL.

Getting the lab SQL script file

- ❑ Download the lab6.sql file as follows
 - 1. login to an arbitrary machine
csl2wkxx.cse.ust.hk where xx=01-40
 - 2. at the command prompt type
csl2wk01:lamngok:235> cd ~
csl2wk01:lamngok:236> wget \
[?http://course.cs.ust.hk/comp3311/labs/lab6.sql](http://course.cs.ust.hk/comp3311/labs/lab6.sql)
- ❑ Log into Oracle database server using SQL*Plus with your password.
- ❑ Execute the script lab6.sql at the prompt
SQL> @lab6.sql

Oracle Indexing 1

- ❑ An index can be created in the Oracle database system to speed up record retrievals.
- ❑ An index is defined on one or more columns.
- ❑ An index occupies extra space and is stored separately from the table.
- ❑ Oracle uses the index only when the index is estimated to improve the performance.

Oracle Indexing 2

- ❑ In Oracle, a unique column known as “**ROWID**” is used to identify records internally.
- ❑ The “**key**” of the index corresponds to the values of the columns on which the index is created.
- ❑ When an index is created, the index entries will hold the values of the “**key**” and the “**ROWID**” of the records.

Oracle Indexing 3

- ❑ Whenever appropriate, the Oracle system uses the index to search for records.
- ❑ Since the index is stored in a balanced tree, data retrieval from the index is faster.
- ❑ The ROWID information obtained from the index will be used by the Oracle system to directly locate the record in the file system.
- ❑ Large tables with large number of records will benefit greatly from the indexes.

Oracle Indexing 4

- The syntax for creating an index

```
CREATE          [UNIQUE]          INDEX  
index_name     ON   table   (column1,  
column2, ...);
```

- The UNIQUE keyword specifies that the column(s) must have unique values.

Oracle Indexing 5

□ An example

```
CREATE UNIQUE INDEX facility_idx  
ON facility (department_id, name);
```

- The above example creates an index on the **department_id** and **name** columns of the table **facility**. The tuples (**department_id, name**) must be unique.

Oracle indexing 6

- ❑ Oracle does not always use the index. An index will not be used in the following scenarios:
 - The SELECT statement does not contain the WHERE clause:
`SELECT * from facility;`
 - The SELECT statement contains the WHERE clause, but the WHERE clause does not refer to the indexed column(s):
`SELECT * FROM facility where no_of_computers=60;`
 - The indexed column(s) is/are modified by some function(s) in the WHERE clause:
`SELECT * FROM facility where
substr(name,1,8)='Computer';`

Oracle indexing 7

- ❑ A function-based index can be created as follows:

```
CREATE INDEX function_idx ON facility  
(substr(name,1,8));
```

- ❑ We can check the names of all the indexes created:

```
SELECT index_name FROM user_indexes;
```

- ❑ We can also drop an index:

```
DROP INDEX index_name;
```

- ❑ To create the function-based index, the user must have the **QUERY REWRITE** system privilege.

Oracle indexing 8

- ❑ When the Oracle database system is creating an index for a table, the table is locked, and will not be available for data manipulations.
- ❑ This could be undesirable in the real-world working environment.
- ❑ Starting from version 8i, Oracle supports online indexing by using the keyword ONLINE:

```
CREATE    UNIQUE INDEX facility_idx ON  
facility (department_id,name) ONLINE;
```

Oracle indexing 9

- ❑ Creating an index could slow down the insertion and deletion operations.
- ❑ If the table is to be updated frequently with the insertion and deletion statements, creating and maintaining the index could be a big overhead to the system.
- ❑ Index is good for tables that are primarily used for querying and tables that do not require to be updated frequently.

Oracle clustering 1

- ❑ Oracle can store tables that are often used together (i.e. in JOIN operations) in the same data blocks. This is known as “clustering”.
- ❑ Clustering reduces unnecessary I/O accesses, thus improves the performance.
- ❑ To place tables in a cluster, the tables need to have a common column.
- ❑ The user needs the “**CREATE CLUSTER**” system privilege in order to create a cluster.

Oracle clustering 2

- ❑ To create a cluster⁺ in Oracle, you need to refer to the following steps:
 - 1. Create a cluster using the **CREATE CLUSTER** command.
 - 2. Create an index on the cluster using the **CREATE INDEX** command. THIS MUST BE DONE BEFORE ANY RECORDS ARE INSERTED INTO THE TABLES.
 - 3. Place the tables into the cluster with the **CLUSTER** option in the **CREATE TABLE** statement.

⁺ Don't worry if you do not have the privilege to create a cluster.

Oracle clustering 3

- ❑ Step 1: creating a cluster
- ❑ The syntax for creating a cluster:
`CREATE CLUSTER clustername (column1 datatype, column2 datatype,...)`
- ❑ An example
`CREATE CLUSTER department_facility (department_id varchar2(4));`
- ❑ The above example create a cluster named department_facility, department_id is known as the cluster key.

⁺ Don't worry if you do not have the privilege to create a cluster.

Oracle clustering 4

- ❑ Step 2: creating an index
- ❑ Oracle requires the cluster to have an index (otherwise record insertion will not be allowed).
- ❑ But Oracle will not build an index automatically for the cluster.
- ❑ The following statement creates an index on the department_facility cluster:

```
CREATE INDEX  
department_facility_idx  
ON CLUSTER department_facility;
```


Oracle clustering 5

- ❑ Step 3: placing tables into a cluster
- ❑ The tables being put into the cluster must have a column that matches with the cluster key.
- ❑ The following statement creates the departments table and put it into the cluster:

```
CREATE TABLE departments  
( department_id      varchar2(4) not null,  
  name varchar2(40),  
  room_number        number(4) )  
CLUSTER department_facility(department_id);
```

Oracle clustering 6

- ❑ The following statement creates the facility table and put it into the cluster:

```
CREATE TABLE facility  
( department_id  varchar2(4) not null,  
  name          varchar2(40),  
  no_of_projectors number(4),  
  no_of_computers number(5))  
CLUSTER department_facility(department_id);
```

- ❑ In Oracle you can not put existing tables into a cluster.
- ❑ The tables are required to be put into the cluster when they are created.
- ❑ To see the names of all of the clusters created we can use the following command:

```
SELECT cluster_name FROM user_clusters;
```

⁺ *Don't worry if you do not have the privilege to create a cluster.*

Oracle clustering 7

- ❑ The following statement drop the cluster we created:

```
DROP CLUSTER department_facility  
[INCLUDING TABLES  
[CASCADE CONSTRAINTS]];
```

- ❑ **INCLUDING TABLES:** drops the cluster as well as tables in that cluster
- ❑ **CASCADE CONSTRAINTS:** drop any referential integrity constraints that refer to primary or unique keys in tables.
- ❑ You cannot uncluster an individual table. Instead you must perform these steps:
 - Create a new table with the same structure and contents as the old one, but with no CLUSTER clause.
 - Drop the old table.
 - Use the RENAME statement to give the new table the name of the old one.

More about PL/SQL 1

- Just a reminder, the basic structure of PL/SQL is as follows:

DECLARE

/* Declarative section: variables, types, and local subprograms. */

BEGIN

/* Executable section: procedural and SQL statements go here. */

/* This is the only section of the block that is required. */

EXCEPTION

/* Exception handling section: error handling statements go here. */

END;

More about PL/SQL 2 (additional flow control statements)

- In addition to the control statements we discussed in the last lab, there are some more control statements in PL/SQL. The following is a list of them (covered statements are in red):

Conditional control statements:

- IF ... THEN ... ELSIF ... ELSE ... END IF;
- CASE ... WHEN ... THEN ... ELSE ... END CASE;

Iterative statements:

- LOOP ... END LOOP;
- WHILE ... LOOP ... END LOOP;
- FOR ... IN ... LOOP ... END LOOP;

More about PL/SQL 3 (exceptions)

☐ Handling exceptions in PL/SQL

■ Predefined exceptions

- ☐ NO_DATA_FOUND, TOO_MANY_ROWS, etc

- ☐ For the complete listing of predefined exceptions, refer to:

http://download.oracle.com/docs/cd/B10501_01/appdev.920/a96624/07_errs.htm

■ User-defined exceptions

- ☐ Defined by the users

- ☐ Raised explicitly by users using the RAISE command:

RAISE <exception name>;

More about PL/SQL 4 (exceptions)

- To use a user-defined exception, we need to:
 - Declare the exception under the “**DECLARE**” section,
 - Raise it (whenever applicable) under the “**BEGIN**” section,
 - Define the codes under the “**EXCEPTION**” section.

More about PL/SQL 5 (exceptions)

- We will show the use of a user-defined exception `cga_too_low` in a piece of PL/SQL code on the next slide.
- The code checks a student with the email address `'lamngok'`.
- If the `cga` value of that student is lower than 10 then a user-defined exception will be raised:
 - A message `'LAM IS LAZY'` will be displayed to the screen.
 - The student's `last_name` will be updated to `'LAZY'`

More about PL/SQL 6 (exceptions)

```
DECLARE
    cga_too_low EXCEPTION;
    cga_of_lam NUMBER(2);
BEGIN
    SELECT cga INTO cga_of_lam
    FROM students WHERE email='lamngok';
    IF (cga_of_lam<10) THEN
        RAISE cga_too_low;
    END IF;

EXCEPTION
    WHEN cga_too_low THEN
        DBMS_OUTPUT.PUT_LINE('LAM IS LAZY!');
        UPDATE students SET last_name='LAZY'
        WHERE email='lamngok';

END;
/
```

+ If SQLPlus does not display the sentence "LAM IS LAZY!", please type 25
"set serveroutput on;" at the SQLPlus prompt and run the PL/SQL block again

Conclusions

- We covered the following topics in this lab:
 - creating an index for an table in Oracle,
 - creating a cluster for tables in Oracle,
 - further PL/SQL knowledge.