

COMP 3511

Operating Systems



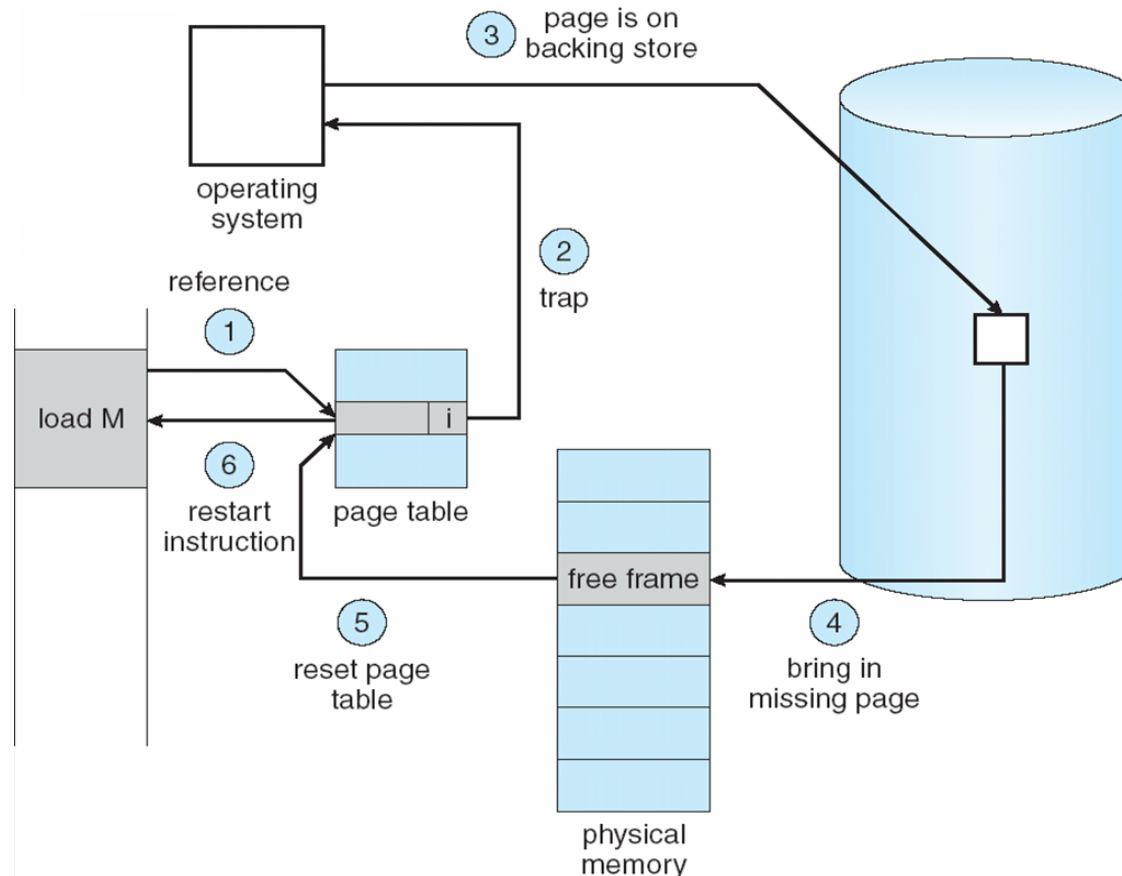
Lab 07

Demand Paging

- Bring a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Page is needed → reference to it
 - Invalid reference → abort
 - Not-in-memory → bring to memory

Steps in Handling a Page Fault

- If there is a reference to a page, first reference to that page will trap to operating system: **page fault**



An example of demand-paged memory

- Assume we have a **demand-paged** memory
 - The page table is held in registers
 - It takes **8 milliseconds** to serve a page fault if an empty page is available or the replaced page is not modified
 - It takes **20 milliseconds** if the replaced page is modified
 - Memory access time is **100 nanoseconds**

An example of demand-paged memory

- Assume that the page to be replaced is modified **70 percent** of the time.
- *What is the maximum acceptable page-fault rate for an effective access time of no more than **200 nanoseconds** ?*

An example of demand-paged memory

$$0.2\mu\text{sec} = (1 - P) \times 0.1\mu\text{sec} + (0.3P) \times 8 \text{ millisec} \\ + (0.7P) \times 20 \text{ millisec}$$

$$0.1 = -0.1P + 2400 P + 14000 P$$

$$0.1 \approx 16,400 P$$

$$P \approx 0.000006$$

Hardware support for demand paging

- For **every memory access operation**, the **page table** needs to be consulted:
 - check whether the corresponding page is resident or not
 - check whether the program has read or write privileges for accessing the page.

Hardware support for demand paging

- These checks would have to be performed in hardware.
- For example, a **TLB** could serve as a cache and improve the performance of the lookup operation.


Page Replacement

- If there is no free frame
- **Page replacement** – find some page in memory, but not really in use, swap it out
 - Replacement algorithm
 - Performance requirement – want an algorithm which will result in minimum number of page faults

Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame inside the memory:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Read the desired page into the (newly) free frame. Update the page and frame tables.
4. Restart the process

Page Replacement



First-In-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults

- 4 frames

1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults

Belady's Anomaly:
more frames →
more page faults

FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2
	0	0	0
		1	1

2	2	4	4	4	0
3	3	3	2	2	2
1	0	0	0	3	3

0	0
1	1
3	2

7	7	7
1	0	0
2	2	1

page frames

Optimal page replacement algorithm

- Replace page that will not be used for longest period of time
- Problem
 - Can't read the future
- Used for measuring how well your algorithm performs


Algorithms for approximating optimal page replacement

- **LRU** (Least Recently Used) algorithm
 - Use the **recent past** as an approximation of the **near future**
 - Replace the page that has not been used for the longest period of time
 - Considered to be good, but hard to implement
 - Few computer systems provide sufficient hardware support for true LRU
 - LRU-approximation: **Reference bits**, **Second chance**

■ Optimal page replacement (9 page faults)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

■ LRU page replacement (12 page faults)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

LRU Approximation Algorithms


■ Reference bit

- Each page is associated a reference bit, initially = 0.
- When page is referenced, the reference bit set to 1.
- We can determine which pages have been used or not used by examining the reference bits and replace the page whose reference bit is 0 (if one exists).
- However, we do not know the *order* of use.


■ Second chance

- When a page is selected, check the reference bit.
- If the value is 0, this page is replaced.
- If the value is 1, set the bit to 0, then move on to select another page (FIFO). This page gets a second chance.

Second chance algorithm



(a)



(b)

Algorithms for approximating optimal page replacement

- Keep a **counter** of the number of references that have been made to each page
 - **LFU** (Least Frequently Used) algorithm
 - an actively used page should have a large reference count
 - replaces page with **smallest count**
 - **MFU** (Most Frequently Used) algorithm
 - the page with the smallest count was probably just brought in and has yet to be used
- The implementation of these algorithms is expensive, and **they do not approximate OPT replacement well**

An example of page-replacement algorithm

MFU: *most frequently used page-replacement algorithm*

VS.

LRU: *least recently used page replacement algorithm*

- (a) Under which situations, the **MFU** generates fewer page faults than the **LRU** ? Please give an example.
- (b) Under what circumstance does the opposite holds ?

Example 1: MFU & LRU

- (a) Consider the sequence in a system that holds **four** pages in memory: **1 2 3 4 4 4 5 1**
 - the **MFU** evicts page **4** while fetching page **5**
 - the **LRU** evicts page **1** while fetching page **5**, then another page fault for fetching page **1** again
- (b) For the sequence “**1 2 3 4 4 4 5 4**,” the **LRU** algorithm makes the right decision

Example 2: LFU & LRU

- Consider the following sequence of memory accesses in a system that can hold **four** pages in memory: 1 1 2 3 4 5 1
- Which of the following page replacement algorithms generates fewer page faults?
 - a) LRU: least recently used page replacement
 - b) LFU: least frequently used page replacement

Example 2: LFU & LRU

- memory accesses: 1 1 2 3 4 5 1
- **four** pages in memory
- The LRU evicts page **1** while fetching page **5**
- The LFU evicts a page other than **1** while fetching page **5**
 - No page fault when page **1** is accessed again
 - fewer page faults