

Chapter 4: Multithreaded Programming



Chapter 4: Multithreaded Programming

- Overview
- Multicore Programming
- Multithreading Models
- Threading Issues
- Operating System Examples



Objectives

- To introduce the notion of a thread—a fundamental unit of CPU utilization that forms the basis of multithreaded computer systems
- To discuss the APIs for the Pthreads, Windows, and Java thread libraries
- To examine issues related to multithreaded programming
- To cover operating system support for threads in Windows and Linux




Motivation

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded



Multithreaded Server Architecture



Benefits

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads with a process share resources of the process, easier than shared memory or message passing
- **Economy** – thread creation is much cheaper than process creation, thread switching also has much lower overhead than context switching (switching to a different process)
- **Scalability** – A process can take advantage of multiprocessor architectures by running multiple threads of the process simultaneously on different processors (CPUs).




Multicore Programming

- **Multicore or multiprocessor** systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency
- Types of parallelism
 - **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
 - **Task parallelism** – distributing threads across cores, each thread performing unique operation




Concurrency vs. Parallelism

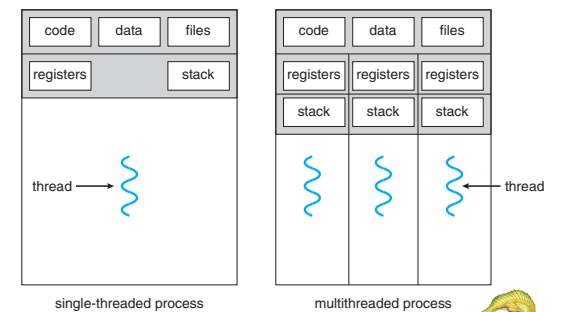
- **Concurrent execution on single-core system:**



- **Parallelism on a multi-core system:**



Single and Multithreaded Processes





Thread State

- Each Thread has a **Thread Control Block (TCB)**
 - Execution State: CPU registers, program counter, pointer to stack
 - Scheduling info: State (more later), priority, CPU time
 - Accounting Info
 - Various Pointers (for implementing scheduling queues)
 - Pointer to enclosing process? (PCB)?
- In Nachos: "thread" is a class that includes the TCB
- OS Keeps track of TCBs in protected memory
 - Array, or Linked List, or ...




Thread State (Cont.)

- State shared by all threads in process/address space
 - Contents of memory (global variables, heap)
 - I/O state (file system, network connections, etc)
- State "private" to each thread
 - Kept in TCB - Thread Control Block
 - CPU registers (including, program counter)
 - Execution stack - what is this?
- Execution Stack
 - Parameters, temporary variables
 - Keep program counters while called procedures are executing



Lifecycle of a Thread

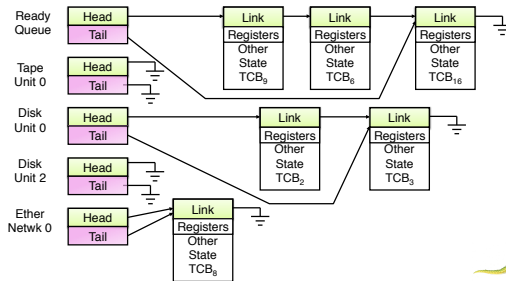


- As a thread executes, it changes state:
 - new:** The thread is being created
 - ready:** The thread is waiting to run
 - running:** Instructions are being executed
 - waiting:** Thread waiting for some event to occur
 - terminated:** The thread has finished execution
- "Active" threads are represented by their TCBs
 - TCBs organized into queues based on their state



Ready Queue And Various I/O Device Queues

- Thread not running → TCB is in some scheduler queue
 - Separate queue for each device/signal/condition
 - Each queue can have a different scheduler policy



Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- S is serial portion
- N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- I.e. If application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1/S$

Serial portion of an application has disproportionate effect on performance gained by adding additional cores



Examples of Multithreaded Programs

- Embedded systems
 - Elevators, Planes, Medical systems, Wristwatches
 - Single Program, concurrent operations
- Most modern OS kernels
 - Internally concurrent to deal with concurrent requests by multiple users
 - But no protection needed within kernel
- Database Servers
 - Access to shared data by many concurrent users
 - Also background utility processing must be done
- Network Servers
 - Concurrent requests from network
 - Again, single program, multiple concurrent operations
 - File server, Web server, and airline reservation systems
- Parallel Programming (More than one physical CPU)
 - Split program into multiple threads for parallelism
 - This is called **Multiprocessing**



User Threads and Kernel Threads

- User threads** - management done by user-level threads library
- Three primary thread libraries:
 - POSIX Pthreads
 - Win32 threads
 - Java threads
- Kernel threads** - Supported by the Kernel
- Examples - virtually all general purpose operating systems, including:
 - Windows
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X




Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many



Many-to-One

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads






One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead

- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later



Operating System Concepts - 9th Edition


4.19

Silberschatz, Galvin and Gagne ©2013



Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows NT/2000 with the *ThreadFiber* package



Operating System Concepts - 9th Edition


4.20

Silberschatz, Galvin and Gagne ©2013



Two-level Model

- Similar to M.M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



Operating System Concepts - 9th Edition

4.21

Silberschatz, Galvin and Gagne ©2013



Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Signal handling
 - Synchronous and asynchronous
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread-local storage



Operating System Concepts - 9th Edition

4.22

Silberschatz, Galvin and Gagne ©2013



Semantics of fork() and exec()

- Does **fork()** duplicate only the calling thread or all threads?
 - Some UNIXes have two versions of fork
- Exec()** usually works as normal – replace the running process including all threads



Operating System Concepts - 9th Edition

4.23

Silberschatz, Galvin and Gagne ©2013



Signal Handling

- Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- A **signal handler** is used to process signals
 - Signal is generated by particular event
 - Signal is delivered to a process
 - Signal is handled by one of two signal handlers:
 - default
 - user-defined
- Every signal has **default handler** that kernel runs when handling signal
 - User-defined signal handler** can override default
 - For single-threaded, signal delivered to process
- Where should a signal be delivered for multi-threaded?
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process



Operating System Concepts - 9th Edition

4.24

Silberschatz, Galvin and Gagne ©2013



Thread Cancellation

- Terminating a thread before it has finished
- Thread to be canceled is **target thread**
- Two general approaches:
 - Asynchronous cancellation** terminates the target thread immediately
 - Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

...

/* cancel the thread */
pthread_cancel(tid);
```



Operating System Concepts - 9th Edition

4.25

Silberschatz, Galvin and Gagne ©2013



Thread-Local Storage

- Thread-local storage (TLS)** allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- Different from local variables
 - Local variables visible only during single function invocation
 - TLS visible across function invocations
- Similar to **static** data
 - TLS is unique to each thread



Operating System Concepts - 9th Edition

4.26

Silberschatz, Galvin and Gagne ©2013



Operating System Examples

- Windows XP Threads
- Linux Thread



Operating System Concepts - 9th Edition

4.27

Silberschatz, Galvin and Gagne ©2013

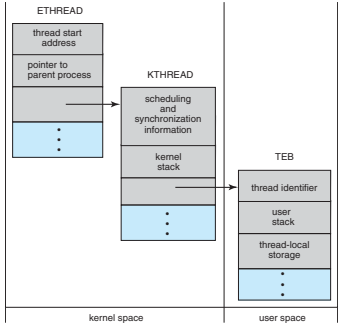


Windows Threads

- Windows implements the Windows API – primary API for Win 98, Win NT, Win 2000, Win XP, and Win 7
- Implements the one-to-one mapping, kernel-level
- Each thread contains
 - A thread id
 - Register set representing state of processor
 - Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- The register set, stacks, and private storage area are known as the **context** of the thread
- The primary data structures of a thread include:
 - ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space
 - KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
 - TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space



Windows XP Threads Data Structures



Linux Threads

- Linux refers to them as **tasks** rather than **threads**
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
 - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- `struct task_struct` points to process data structures (shared or unique)

