

Spring 2015 - COMP3511 Final Exam Review



Outline

- Deadlock and Banker Algorithm
- Paging and Segmentation
- Page Replacement Algorithms and Working-set Model
- File Allocation
- Disk Scheduling

Operating System Concepts

Review.2



Deadlock Characterization

- Deadlock can arise if **four** conditions hold simultaneously
- Mutual exclusion**
 - only one process at a time can use a resource.
- Hold and wait**
 - a process holding at least one resource is waiting to acquire additional resources held by other processes.
- No preemption**
 - a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- Circular wait**
 - there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Operating System Concepts

Review.3



Resource-Allocation Graph

- A set of vertices V and a set of edges E .
- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- Each resource type R_i has W_i instances.
- Each process utilizes a resource as follows: **request, use, release**
- Request edge** – directed edge $P_i \rightarrow R_j$
- Assignment edge** – directed edge $R_j \rightarrow P_i$

Operating System Concepts

Review.4



Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in safe state if there exists a safe sequence of all processes:
 - Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe if for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on

Operating System Concepts

Review.5



Banker's Algorithm

- Each resource can have multiple instances.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

Let n = number of processes, and m = number of resources types.

- Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available.
- Max:** $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .
- Allocation:** $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j .
- Need:** $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i, j] = Max[i, j] - Allocation[i, j].$$

Operating System Concepts

Review.6



Safety Algorithm

- Let **Work** and **Finish** be vectors of length m and n , respectively.
Initialize:
 $Work = Available$
 $Finish[i] = false$ for $i = 1, 3, \dots, n$.
- Find and i such that both:
 - $Finish[i] = false$
 - $Need_i \leq Work$If no such i exists, go to step 4
- $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2
- If $Finish[i] == true$ for all i , then the system is in a safe state

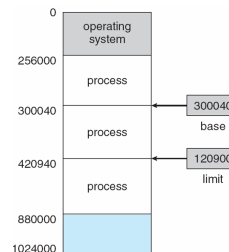
Operating System Concepts

Review.7



Base and Limit Registers

- Two special registers, **base** and **limit** are used to prevent user from straying outside the designated area
- During context switch, OS loads new base and limit register from TCB
- User is NOT allowed to change the base and limit registers (privileged instructions)



Operating System Concepts

Review.8



Contiguous memory allocation

- Each process is contained in a **single contiguous section of memory**
 - Hole:** block of available memory
 - holes of various size are scattered throughout memory
- Operating system maintains information about
 - a) allocated partitions
 - b) free partitions (hole)

Operating System Concepts

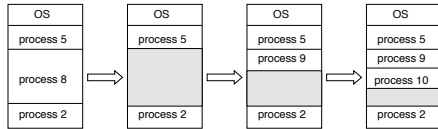
Review.9





Contiguous Memory Allocation

- Each process is contained in a **single contiguous section of memory**
 - Degree of multiprogramming limited by number of partitions
 - Variable-partition** sizes for efficiency (sized to a given process' needs)
 - Hole** – block of available memory; holes of various size are scattered throughout memory
 - Operating system maintains information about:
 - allocated partitions
 - free partitions (hole)



Operating System Concepts

Review.10



Paging

- Physical address space of a process can be **noncontiguous**
 - Divide physical memory into fixed-sized blocks called **frames**,
 - Divide logical memory into blocks of same size called **pages**.
 - Keep track of all free frames
- Set up a **page table** to translate logical to physical addresses

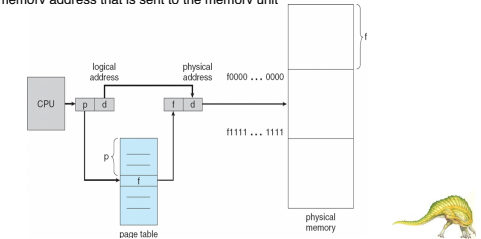
Operating System Concepts

Review.11



Address Translation

- Address generated by CPU is divided into:
 - Page number (p)** – used as an index into a **page table** which contains base address of each page in physical memory
 - Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit



Operating System Concepts

Review.12



Page Table Implementation

- Implementation of Page Table
 - Page table** is kept in **main memory**
 - Page-table base register (PTBR)** points to the page table
 - Page-table length register (PRLR)** indicates size of the page table
 - In this scheme **every data/instruction access requires two memory accesses**.
 - One for the page table and one for the data/instruction

Operating System Concepts

Review.13



TLB

- The two memory access problem can be solved by using **TLB (translation look-aside buffer)**
 - a special, small, fast-lookup hardware cache
 - each entry in the TLB consists of a key (or tag) and a value
 - page number is presented to the TLB, if found, its frame number is immediately available to access memory
 - fast but expensive

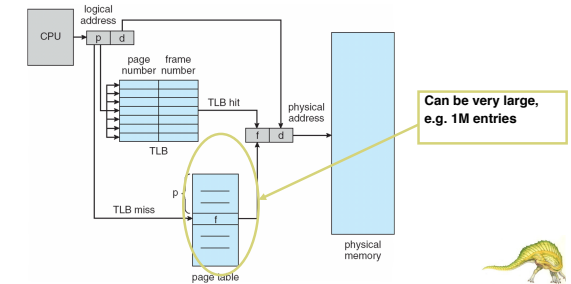
Operating System Concepts

Review.14



Paging Hardware With TLB

- The two memory access problem can be solved by using **TLB (translation look-aside buffer)**



Operating System Concepts

Review.15



TLB miss and Hit ratio

- TLB miss:**
 - If the page number is not in the TLB, a memory reference to the page table must be made
- Hit ratio:**
 - percentage of times that a page number is found in the TLB.
- For example:
 - Assume TLB search takes 20ns; memory access takes 100ns
 - TLB hit → 1 memory access; TLB miss → 2 memory accesses

Operating System Concepts

Review.16



Effective Access Time (EAT)

- If Hit ratio = 80%
 - $EAT = (20 + 100) * 0.8 + (20 + 200) * 0.2 = 140ns$
- If Hit ratio = 98%
 - $EAT = (20 + 100) * 0.98 + (20 + 200) * 0.02 = 122ns$

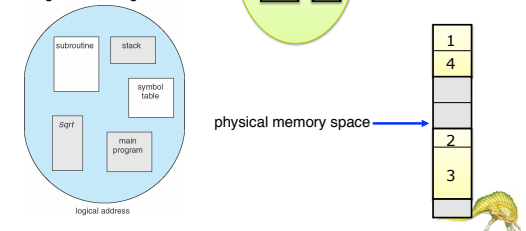
Operating System Concepts

Review.17



Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of **segments of different sizes**
- A segment is a logical unit



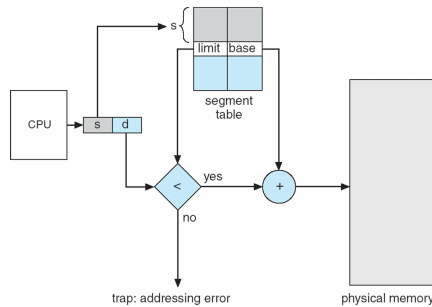
Operating System Concepts

Review.18





Address Translation

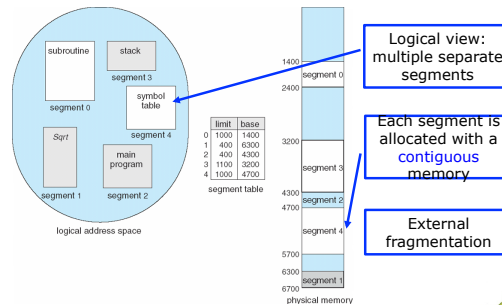


Operating System Concepts

Review.19



Example of Segmentation



Operating System Concepts

Review.20



Paging-Segmentation Combination

- Segmentation and Paging are often combined in order to improve upon each other
- Segmented paging is helpful when the page table becomes very large
 - e.g., a large contiguous section of the page table that is unused can be collapsed into a single segment table entry with a page table address of zero

Operating System Concepts

Review.21



Motivation of virtual memory

- Should an entire process be in memory before it can execute?
 - In fact, real programs show us that, in many cases, the entire program is not needed
 - Even in those cases where the entire program is needed, it may not all be needed at the same time
 - More programs could run concurrently, increasing CPU utilization and throughput
 - Less I/O would be needed to load or swap each user program into memory, so each user program would run faster
 - Allow processes to share files easily and to implement shared memory

Operating System Concepts

Review.22



Page Replacement

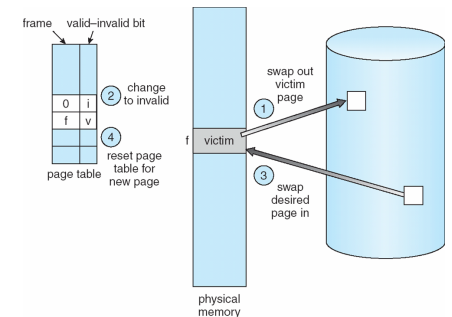
- If there is no free frame
- Page replacement – find some page in memory, but not really in use, swap it out
 - Replacement algorithm – local replacement
 - performance – want an algorithm which will result in minimum number of page faults
 - Same page may be brought into and out of memory several times

Operating System Concepts

Review.23



Page Replacement



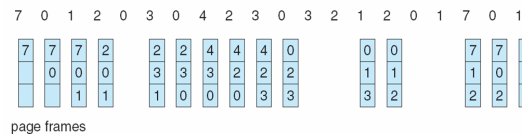
Operating System Concepts

Review.24



FIFO Page Replacement

reference string



Operating System Concepts

Review.25



Algorithms for approximating optimal page replacement

- LRU (Least Recently Used) algorithm
 - Use the recent past as an approximation of the near future
 - Replace the page that has not been used for the longest period of time, to approximate optimal replacement or OPT algorithm, which selects a page that will not be used in the longest amount of time in the future.
 - Considered to be good, but how to implement
 - Few computer systems provide sufficient hardware support for true LRU
 - LRU-approximation: Reference bits, Second chance

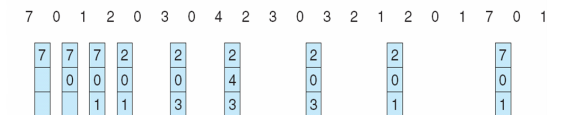
Operating System Concepts

Review.26



Optimal page replacement (9 page faults)

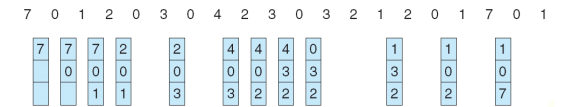
reference string



page frames

LRU page replacement (12 page faults)

reference string



page frames

Operating System Concepts

Review.27

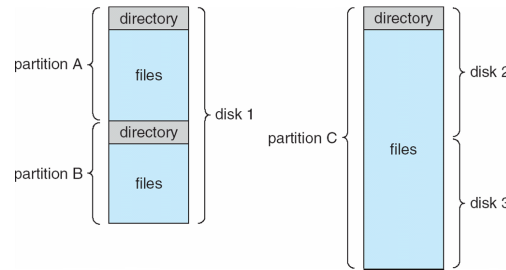


Working-Set Model

- Working-Set model is based on the locality of memory access
- Δ = **working-set window** = a fixed number of page references
Example: 10,000 instructions
- WSS** (**working set** of Process P) = total number of pages referenced in the most recent Δ (varies in time)
 - if Δ too small will not encompass entire locality
 - if Δ too large will encompass several localities
 - if $\Delta = \infty \Rightarrow$ will encompass entire program
- $D = \sum WSS_i$ = total demand for frames (by all processes)
 - if $D > m \Rightarrow$ Thrashing (m is the available frames)
 - Policy if $D > m$, then suspend one of the processes; the process pages are swapped out, and its frames are re-allocated to other processes. The suspended process can be re-started later



A Typical File-system Organization

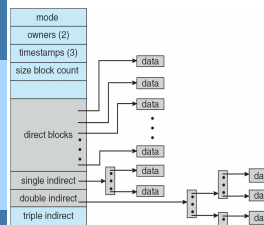


Allocation Methods

- An allocation method refers to how disk blocks are allocated for files – Objectives:
 - Maximize sequential performance
 - Easy random access to file
 - Easy management of file (growth, truncation, and etc)
- Contiguous allocation
- Linked allocation
- Indexed allocation



Combined Scheme: UNIX (4K bytes per block)



- Multi-level index file, key idea:
 - Efficient for small files, still allow large files
- File header format are:
 - First 10 pointers are to data blocks
 - Pointer 11 points to "indirect block" containing 256 block ptrs
 - Pointer 12 points to "doubly indirect block" containing 256 indirect block pointers for total of 64K blocks
 - Pointer 13 points to a triply indirect block (16M blocks)
- Pointers get filled in dynamically



Free-Space Management

- Bit vector (n blocks)

0 1 2 ... n-1

0
1
2
...
n-1

$$\text{bit}[i] = \begin{cases} 1 & \Rightarrow \text{block}[i] \text{ free} \\ 0 & \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$
- Linked list (free list) (previous block contains a pointer to the next free block)
 - Cannot get contiguous space easily
 - No waste of space
- Grouping (stores the addresses of n free blocks in the first free block)
 - Counting
 - Several contiguous blocks may be allocated and freed simultaneously <first free block, number of free contiguous blocks>



Disk Scheduling

- The operating system is responsible for using hardware efficiently — for the magnetic disk drives, this means having a fast access time and disk bandwidth.
- Access time has two major components
 - Seek time** is the time for the disk are to move the heads to the cylinder (tracks) containing the desired sector.
 - Rotational latency** is the additional time waiting for the disk to rotate the desired sector to the disk head.
- Minimize seek time
- Seek time \approx seek distance
- The disk **bandwidth** is the total number of bytes transferred, divided by the total time between the first request for service and the completion of the last transfer.



Disk Scheduling

- When a process needs I/O to or from a disk, it issues a system call to the OS containing the following pieces of information
 - Whether the operation is input or output
 - What the disk address for the transfer is
 - What memory address for the transfer is
 - What the number of sectors to be transferred is
- The question is, when one request is completed, the OS needs to choose which pending requests to service next? How does the OS make this choice?
 - We need *disk scheduling algorithms*
 - FCFS, SSTF, SCAN and LOOK, C-SCAN and C-LOOK

