

## Chapter 3 Transport Layer



Transport Layer 3-1

## Chapter 3: Transport Layer

### our goals:

- ❖ understand principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- ❖ learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport
  - TCP congestion control

Transport Layer 3-2


## Chapter 3 outline

- 3.1 transport-layer services
- 3.2 multiplexing and demultiplexing
- 3.3 connectionless transport: UDP
- 3.4 principles of reliable data transfer
- 3.5 connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- 3.6 principles of congestion control
- 3.7 TCP congestion control

Transport Layer 3-3

## Transport services and protocols

- ❖ provide *logical communication* between app processes running on different hosts
- ❖ transport protocols run in end systems
  - send side: breaks app messages into *segments*, passes to network layer
  - rcv side: reassembles segments into messages, passes to app layer
- ❖ more than one transport protocol available to apps
  - Internet: TCP and UDP



Transport Layer 3-4

## Transport vs. network layer

- ❖ *network layer*: logical communication between hosts
  - ❖ *transport layer*: logical communication between processes
    - relies on, enhances, network layer services
- household analogy:*


12 kids in Ann's house sending letters to 12 kids in Bill's house:

  - ❖ hosts = houses
  - ❖ processes = kids
  - ❖ app messages = letters in envelopes
  - ❖ transport protocol = Ann and Bill who demux to in-house siblings
  - ❖ network-layer protocol = postal service

Transport Layer 3-5

## Internet transport-layer protocols

- ❖ reliable, in-order delivery (TCP)
  - congestion control
  - flow control
  - connection setup
- ❖ unreliable, unordered delivery: UDP
  - no-frills extension of "best-effort" IP
- ❖ services not available:
  - delay guarantees
  - bandwidth guarantees



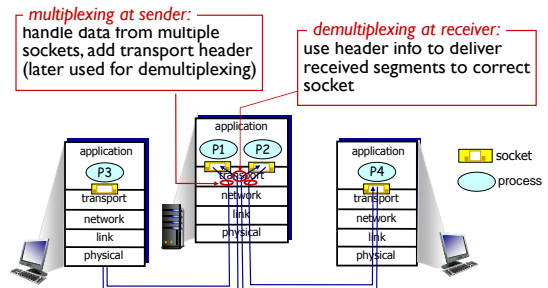
Transport Layer 3-6

## Chapter 3 outline

- 3.1 transport-layer services
- 3.2 multiplexing and demultiplexing
- 3.3 connectionless transport: UDP
- 3.4 principles of reliable data transfer
- 3.5 connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- 3.6 principles of congestion control
- 3.7 TCP congestion control

Transport Layer 3-7


## Multiplexing/demultiplexing



Transport Layer 3-8

## How demultiplexing works

- ❖ host receives IP datagrams
  - each datagram has source IP address, destination IP address
  - each datagram carries one transport-layer segment
  - each segment has source, destination port number
- ❖ host uses **IP addresses & port numbers** to direct segment to appropriate socket



Transport Layer 3-9

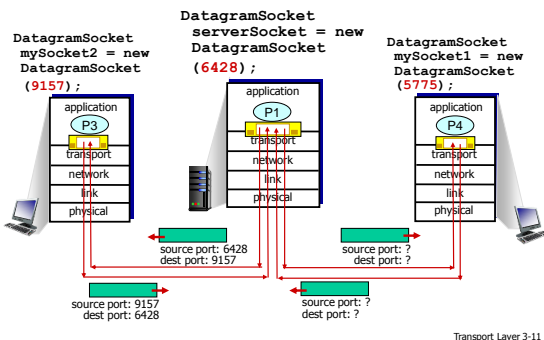
## Connectionless demultiplexing

- ❖ **recall:** created socket has host-local port #:
 

```
DatagramSocket mySocket1 = new DatagramSocket(12534);
```
  - ❖ **recall:** when creating datagram to send into UDP socket, must specify
    - destination IP address
    - destination port #
  - ❖ when host receives UDP segment:
    - checks destination port # in segment
    - directs UDP segment to socket with that port #
- IP datagrams with **same dest. port #**, but different source IP addresses and/or source port numbers will be directed to **same socket** at dest

Transport Layer 3-10

## Connectionless demux: example



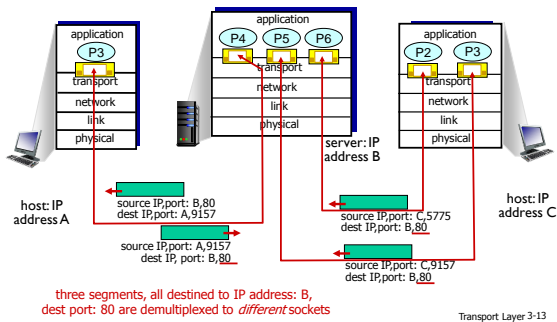
Transport Layer 3-11

## Connection-oriented demux

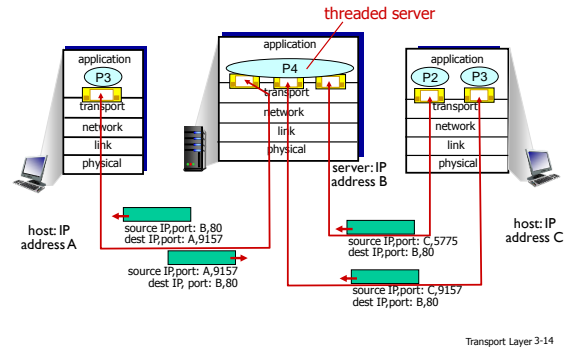
- ❖ TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- ❖ demux: receiver uses all four values to direct segment to appropriate socket
- ❖ server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- ❖ web servers have different sockets for each connecting client
  - non-persistent HTTP will have different socket for each request

Transport Layer 3-12

## Connection-oriented demux: example



## Connection-oriented demux: example



## Chapter 3 outline

- 3.1 transport-layer services
- 3.2 multiplexing and demultiplexing
- 3.3 **connectionless transport: UDP**
- 3.4 principles of reliable data transfer
- 3.5 connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- 3.6 principles of congestion control
- 3.7 TCP congestion control

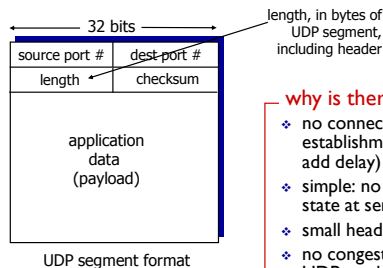
Transport Layer 3-15

## UDP: User Datagram Protocol [RFC 768]

- ❖ “no frills,” “bare bones” Internet transport protocol
- ❖ “best effort” service, UDP segments may be:
  - lost
  - delivered out-of-order to app
- ❖ **connectionless:**
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others
- ❖ UDP use:
  - streaming multimedia apps (loss tolerant, rate sensitive)
  - DNS
  - SNMP
- ❖ reliable transfer over UDP:
  - add reliability at application layer
  - application-specific error recovery!

Transport Layer 3-16

## UDP: segment header



### why is there a UDP?

- ❖ no connection establishment (which can add delay)
- ❖ simple: no connection state at sender, receiver
- ❖ small header size
- ❖ no congestion control: UDP can blast away as fast as desired

Transport Layer 3-17

## UDP checksum

**Goal:** detect “errors” (e.g., flipped bits) in transmitted segment

### sender:

- ❖ treat segment contents, including header fields, as sequence of 16-bit integers
- ❖ checksum: addition (one's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

### receiver:

- ❖ compute checksum of received segment
- ❖ check if computed checksum equals checksum field value:
  - NO - error detected
  - YES - no error detected. But maybe errors nonetheless? More later

Transport Layer 3-18

## Internet checksum: example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
wraparound	①	1	0	1	1	1	0	1	1	1	0	1	1	0	1	1
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	0	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

Transport Layer 3-19


## Chapter 3 outline

- 3.1 transport-layer services
- 3.2 multiplexing and demultiplexing
- 3.3 connectionless transport: UDP
- 3.4 principles of reliable data transfer
- 3.5 connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- 3.6 principles of congestion control
- 3.7 TCP congestion control

Transport Layer 3-20

## Principles of reliable data transfer

- ❖ important in application, transport, link layers
  - top-10 list of important networking topics!




- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Transport Layer 3-21

## Principles of reliable data transfer

- ❖ important in application, transport, link layers
  - top-10 list of important networking topics!




- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Transport Layer 3-22

## Principles of reliable data transfer

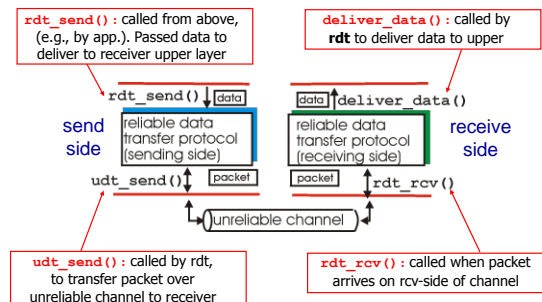
- ❖ important in application, transport, link layers
  - top-10 list of important networking topics!



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Transport Layer 3-23

## Reliable data transfer: getting started




Transport Layer 3-24

## Reliable data transfer: getting started


we'll:

- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
  - but control info will flow on both directions!
- use finite state machines (FSM) to specify sender, receiver



## rdt1.0: reliable transfer over a reliable channel

- underlying channel perfectly reliable
  - no bit errors
  - no loss of packets
- separate FSMs for sender, receiver:
  - sender sends data into underlying channel
  - receiver reads data from underlying channel



## rdt2.0: channel with bit errors

- underlying channel may flip bits in packet
  - checksum to detect bit errors
- the question: how to recover from errors:

*How do humans recover from "errors" during conversation?*

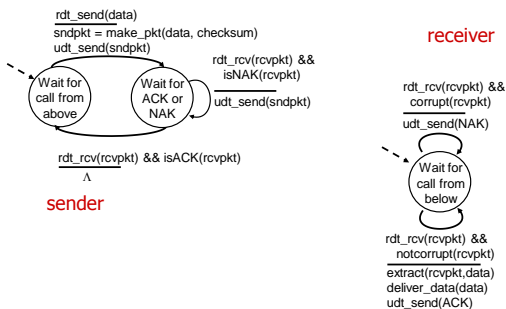
## rdt2.0: channel with bit errors

- underlying channel may flip bits in packet
  - checksum to detect bit errors
- the question: how to recover from errors:
  - acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
  - negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors
  - sender retransmits pkt on receipt of NAK
- new mechanisms in rdt2.0 (beyond rdt1.0):
  - error detection
  - feedback: control msgs (ACK, NAK) from receiver to sender

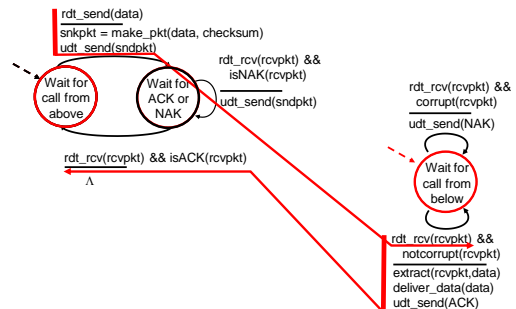
Transport Layer 3-27

Transport Layer 3-28

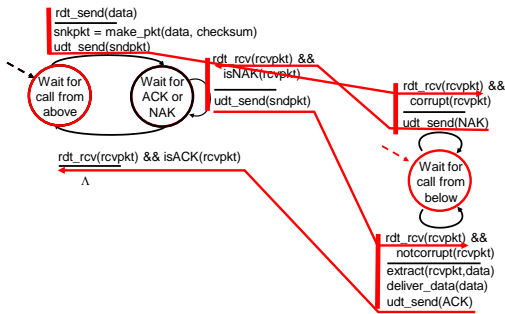
## rdt2.0: FSM specification



## rdt2.0: operation with no errors



## rdt2.0: error scenario



Transport Layer 3-31

## rdt2.0 has a fatal flaw!

### what happens if ACK/NAK corrupted?

- ❖ sender doesn't know what happened at receiver!
- ❖ can't just retransmit: possible duplicate

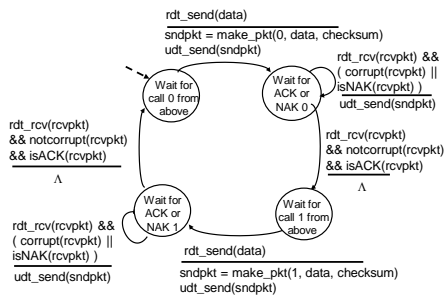
### handling duplicates:

- ❖ sender retransmits current pkt if ACK/NAK corrupted
- ❖ sender adds *sequence number* to each pkt
- ❖ receiver discards (doesn't deliver up) duplicate pkt

**stop and wait**  
sender sends one packet, then waits for receiver response

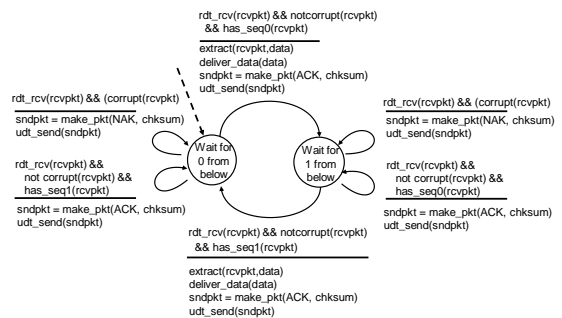
Transport Layer 3-32

## rdt2.1: sender, handles garbled ACK/NAKs



Transport Layer 3-33

## rdt2.1: receiver, handles garbled ACK/NAKs



Transport Layer 3-34

## rdt2.1: discussion

### sender:

- ❖ seq # added to pkt
- ❖ two seq. #'s (0,1) will suffice. Why?
- ❖ must check if received ACK/NAK corrupted
- ❖ twice as many states
  - state must "remember" whether "expected" pkt should have seq # of 0 or 1

### receiver:

- ❖ must check if received packet is duplicate
  - state indicates whether 0 or 1 is expected pkt seq #
- ❖ note: receiver can *not* know if its last ACK/NAK received OK at sender

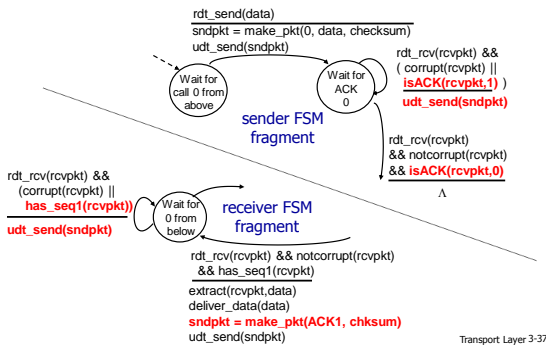
Transport Layer 3-35

## rdt2.2: a NAK-free protocol

- ❖ same functionality as rdt2.1, using ACKs only
- ❖ instead of NAK, receiver sends ACK for last pkt received OK
  - receiver must *explicitly* include seq # of pkt being ACKed
- ❖ duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

Transport Layer 3-36

## rdt2.2: sender, receiver fragments



## rdt3.0: channels with errors and loss

**new assumption:**  
underlying channel can also lose packets (data, ACKs)

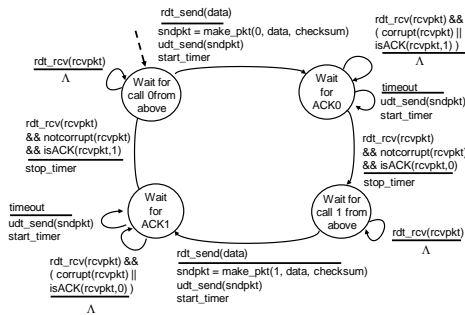
- checksum, seq. #, ACKs, retransmissions will be of help ... but not enough

**approach:** sender waits "reasonable" amount of time for ACK

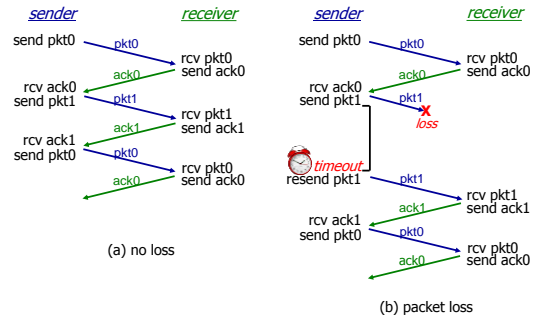
- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
  - retransmission will be duplicate, but seq. #'s already handles this
  - receiver must specify seq # of pkt being ACKed
- requires countdown timer

Transport Layer 3-38

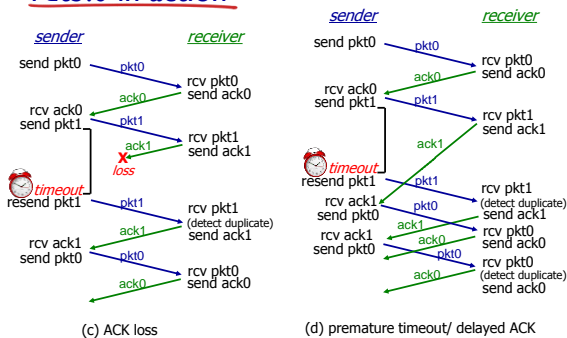
## rdt3.0 sender



## rdt3.0 in action



## rdt3.0 in action



## Performance of rdt3.0

- rdt3.0 is correct, but performance stinks
- e.g.: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microsecs}$$


- $U_{sender}$ : utilization – fraction of time sender busy sending

$$U_{sender} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$

- if RTT=30 msec, 1 KB pkt every 30 msec: 33kB/sec throughput over 1 Gbps link
- network protocol limits use of physical resources!

Transport Layer 3-42

## rdt3.0: stop-and-wait operation




Transport Layer 3-43

## Pipelined protocols

**pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts


- range of sequence numbers must be increased
- buffering at sender and/or receiver



- two generic forms of pipelined protocols: **go-Back-N**, **selective repeat**

Transport Layer 3-44

## Pipelining: increased utilization



Transport Layer 3-45

## Pipelined protocols: overview

### Go-back-N:

- sender can have up to N unacked packets in pipeline
- receiver only sends **cumulative ack**
  - doesn't ack packet if there's a gap
- sender has timer for oldest unacked packet
  - when timer expires, retransmit **all** unacked packets

### Selective Repeat:

- sender can have up to N unack'ed packets in pipeline
- rcvr: sends **individual ack** for each packet
- sender maintains timer for each unacked packet
  - when timer expires, retransmit only that unacked packet

Transport Layer 3-46

## Go-Back-N: sender

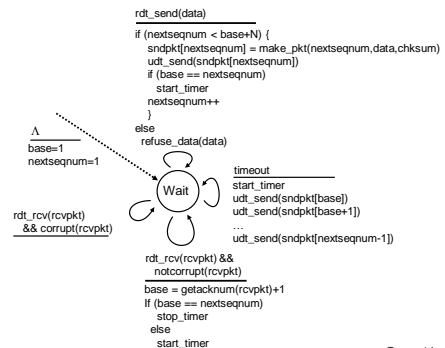
- k-bit seq # in pkt header
- “window” of up to N, consecutive unack'ed pkts allowed



- ACK(n): ACKs all pkts up to, including seq # n - “**cumulative ACK**”
  - may receive duplicate ACKs (see receiver)
- timer for oldest in-flight pkt
- timeout(n): retransmit packet n and all higher seq # pkts in window

Transport Layer 3-47


## GBN: sender extended FSM



Transport Layer 3-48



## GBN: receiver extended FSM



ACK-only: always send ACK for correctly-received pkt with highest **in-order** seq #

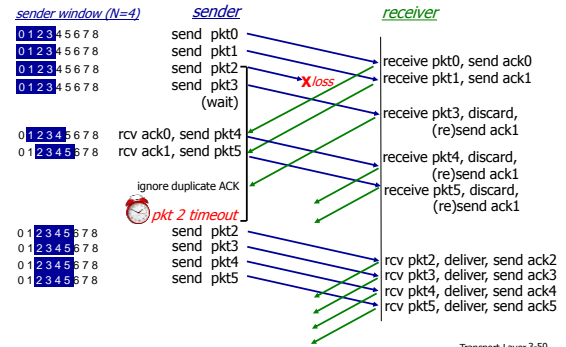
- may generate duplicate ACKs
- need only remember **expectedseqnum**

❖ out-of-order pkt:

- discard (don't buffer): **no receiver buffering!**
- re-ACK pkt with highest in-order seq #

Transport Layer 3-49

## GBN in action



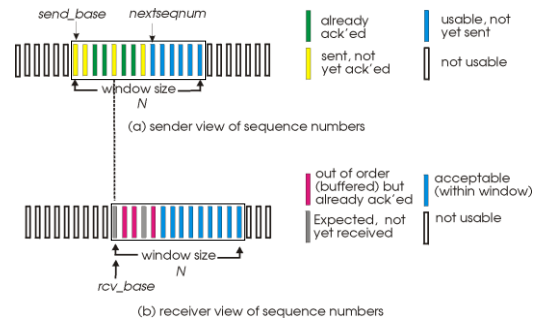
Transport Layer 3-50

## Selective repeat

- ❖ receiver *individually* acknowledges all correctly received pkts
  - buffers pkts, as needed, for eventual in-order delivery to upper layer
- ❖ sender only resends pkts for which ACK not received
  - sender timer for each unACKed pkt
- ❖ sender window
  - N consecutive seq #'s
  - limits seq #'s of sent, unACKed pkts

Transport Layer 3-51

## Selective repeat: sender, receiver windows



Transport Layer 3-52

## Selective repeat

sender

data from above:

- ❖ if next available seq # in window, send pkt

timeout(n):

- ❖ resend pkt n, restart timer

ACK(n) in [sendbase, sendbase+N]:

- ❖ mark pkt n as received
- ❖ if n smallest unACKed pkt, advance window base to next unACKed seq #

receiver

pkt n in [rcvbase, rcvbase+N-1]

- ❖ send ACK(n)
- ❖ out-of-order: buffer
- ❖ in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

pkt n in [rcvbase-N, rcvbase-1]

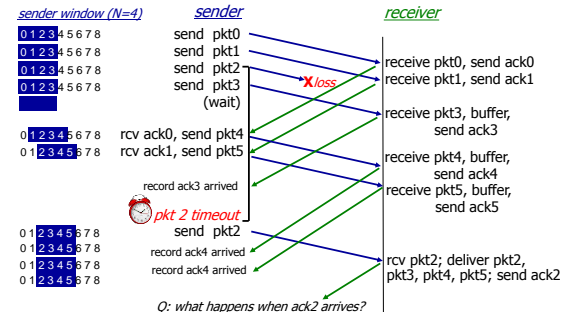
❖ ACK(n)

otherwise:

- ❖ ignore

Transport Layer 3-53

## Selective repeat in action



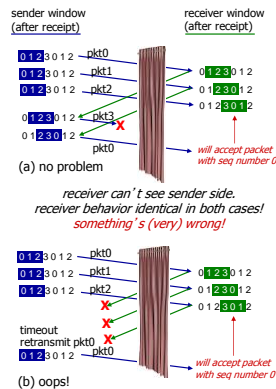
Transport Layer 3-54

## Selective repeat: dilemma

example:

- ❖ seq #'s: 0, 1, 2, 3
- ❖ window size=3
- ❖ receiver sees no difference in two scenarios!
- ❖ duplicate data accepted as new in (b)

Q: what relationship between seq # size and window size to avoid problem in (b)?



## Chapter 3 outline

### 3.1 transport-layer services

### 3.2 multiplexing and demultiplexing

### 3.3 connectionless transport: UDP

### 3.4 principles of reliable data transfer

### 3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

### 3.6 principles of congestion control

### 3.7 TCP congestion control

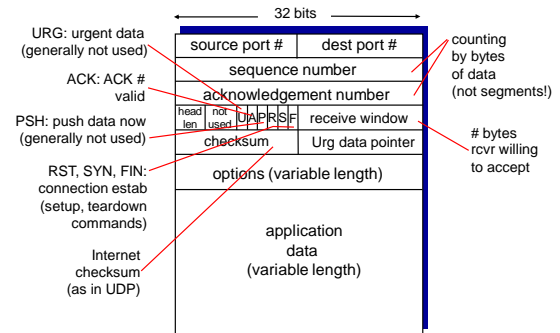
Transport Layer 3-56

## TCP: Overview RFCs: 793, 1122, 1323, 2018, 2581

- ❖ **point-to-point:**
  - one sender, one receiver
- ❖ **reliable, in-order byte stream:**
  - no "message boundaries"
- ❖ **pipelined:**
  - TCP congestion and flow control set window size
- ❖ **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- ❖ **connection-oriented:**
  - handshaking (exchange of control msgs) initiates sender, receiver state before data exchange
- ❖ **flow controlled:**
  - sender will not overwhelm receiver

Transport Layer 3-57

## TCP segment structure



## TCP seq. numbers, ACKs

### sequence numbers:

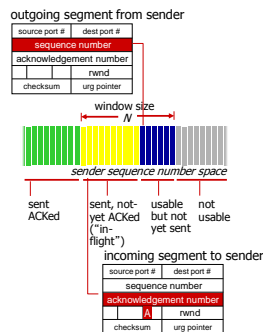
- byte stream "number" of first byte in segment's data

### acknowledgements:

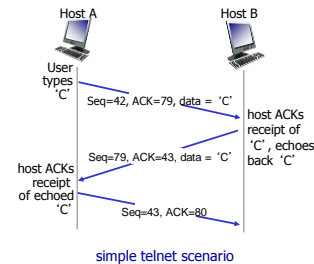
- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn't say, - up to implementor



## TCP seq. numbers, ACKs



Transport Layer 3-60

## TCP round trip time, timeout

Q: how to set TCP timeout value?

- ❖ longer than RTT
  - but RTT varies
- ❖ too short: premature timeout, unnecessary retransmissions
- ❖ too long: slow reaction to segment loss

Q: how to estimate RTT?


- ❖ **SampleRTT**: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- ❖ **SampleRTT** will vary, want estimated RTT "smoother"
  - average several recent measurements, not just current **SampleRTT**

Transport Layer 3-61

## TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- ❖ exponential weighted moving average
- ❖ influence of past sample decreases exponentially fast
- ❖ typical value:  $\alpha = 0.125$




Transport Layer 3-62

## TCP round trip time, timeout

- ❖ **timeout interval**: **EstimatedRTT** plus "safety margin"
  - large variation in **EstimatedRTT** → larger safety margin
- ❖ estimate **SampleRTT** deviation from **EstimatedRTT**:
 
$$\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$


↑ estimated RTT      ↑ "safety margin"

Transport Layer 3-63

## Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Transport Layer 3-64

## TCP reliable data transfer

- ❖ TCP creates rdt service on top of IP's unreliable service

- pipelined segments
- cumulative acks
- single retransmission timer

- ❖ retransmissions triggered by:

- timeout events
- duplicate acks

let's initially consider simplified TCP sender:

- ignore duplicate acks
- ignore flow control, congestion control

Transport Layer 3-65

## TCP sender events:

*data rcvd from app:*

- ❖ create segment with seq #
- ❖ seq # is byte-stream number of first data byte in segment
- ❖ start timer if not already running
  - think of timer as for oldest unacked segment
  - expiration interval: **TimeoutInterval**

*timeout:*

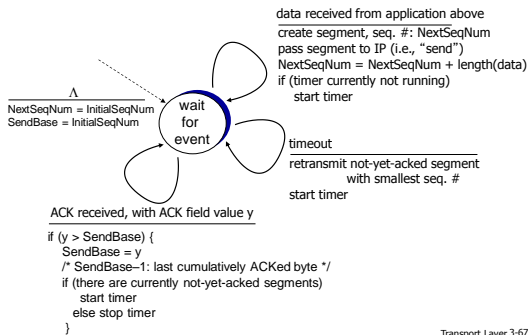
- ❖ retransmit segment that caused timeout
- ❖ restart timer

*ack rcvd:*

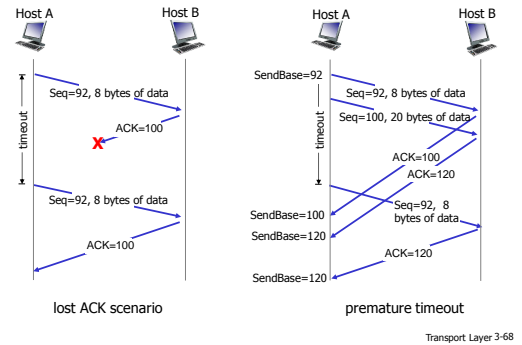
- ❖ if ack acknowledges previously unacked segments
  - update what is known to be ACKed
  - start timer if there are still unacked segments

Transport Layer 3-66


## TCP sender (simplified)



## TCP: retransmission scenarios



## TCP: retransmission scenarios



## TCP ACK generation [RFC 1122, RFC 2581]

event at receiver	TCP receiver action
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expected seq. #. Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

Transport Layer 3-70

## TCP fast retransmit


- ❖ time-out period often relatively long:
  - long delay before resending lost packet
- ❖ detect lost segments via duplicate ACKs.
  - sender often sends many segments back-to-back
  - if segment is lost, there will likely be many duplicate ACKs.

**TCP fast retransmit**  
if sender receives 3 ACKs for same data ("triple duplicate ACKs"), resend unacked segment with smallest seq #

- likely that unacked segment lost, so don't wait for timeout

Transport Layer 3-71

## TCP fast retransmit

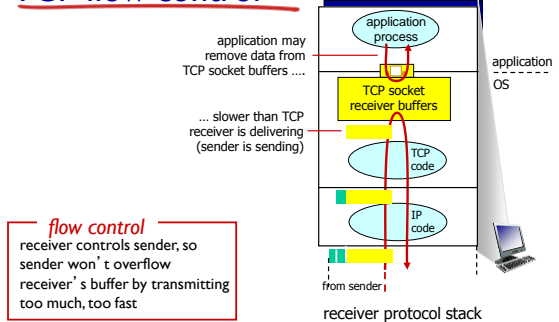


## Chapter 3 outline

- 3.1 transport-layer services
- 3.2 multiplexing and demultiplexing
- 3.3 connectionless transport: UDP
- 3.4 principles of reliable data transfer
- 3.5 connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- 3.6 principles of congestion control
- 3.7 TCP congestion control

Transport Layer 3-73


## TCP flow control



Transport Layer 3-74

## TCP flow control

- ❖ receiver "advertises" free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems autoadjust **RcvBuffer**
- ❖ sender limits amount of unacked ("in-flight") data to receiver's **rwnd** value
- ❖ guarantees receive buffer will not overflow



Transport Layer 3-75

## Chapter 3 outline


- 3.1 transport-layer services
- 3.2 multiplexing and demultiplexing
- 3.3 connectionless transport: UDP
- 3.4 principles of reliable data transfer
- 3.5 connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- 3.6 principles of congestion control
- 3.7 TCP congestion control

Transport Layer 3-76

## Connection Management


before exchanging data, sender/receiver "handshake":

- ❖ agree to establish connection (each knowing the other willing to establish connection)
- ❖ agree on connection parameters



## Agreeing to establish a connection

2-way handshake:



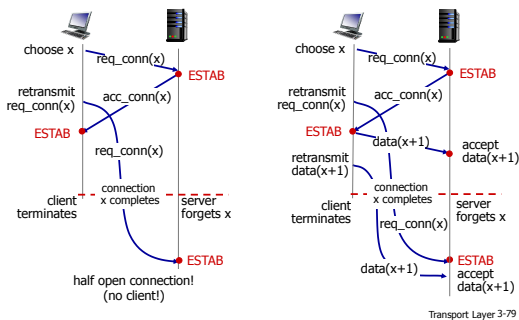
**Q:** will 2-way handshake always work in network?

- ❖ variable delays
- ❖ retransmitted messages (e.g. req\_conn(x)) due to message loss
- ❖ message reordering
- ❖ can't "see" other side

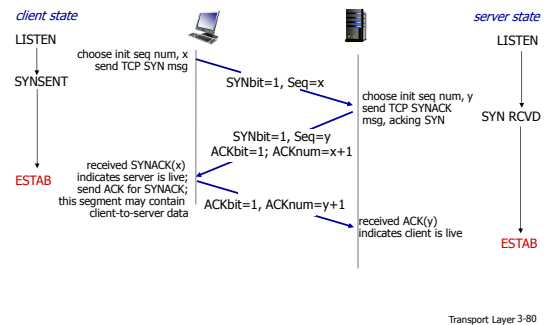
Transport Layer 3-78

## Agreeing to establish a connection

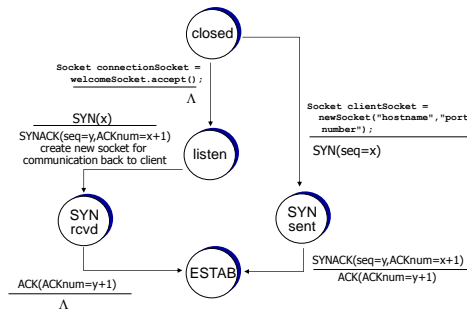
2-way handshake failure scenarios:



## TCP 3-way handshake



## TCP 3-way handshake: FSM




## TCP: closing a connection

- ❖ client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- ❖ respond to received FIN with ACK
  - on receiving FIN, ACK can be combined with own FIN
- ❖ simultaneous FIN exchanges can be handled

Transport Layer 3-82

## TCP: closing a connection



## Chapter 3 outline

- 3.1 transport-layer services
- 3.2 multiplexing and demultiplexing
- 3.3 connectionless transport: UDP
- 3.4 principles of reliable data transfer
- 3.5 connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- 3.6 principles of congestion control
- 3.7 TCP congestion control

Transport Layer 3-84

## Principles of congestion control


### congestion:

- informally: “too many sources sending too much data too fast for *network* to handle”
- different from flow control!
- manifestations:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)
- a top-10 problem!

Transport Layer 3-85

## Causes/costs of congestion: scenario 1

- two senders, two receivers
- one router, infinite buffers
- output link capacity:  $R$
- no retransmission




- maximum per-connection throughput:  $R/2$
- large delays as arrival rate,  $\lambda_{in}$ , approaches capacity

Transport Layer 3-86

## Causes/costs of congestion: scenario 2

- one router, *finite* buffers
- sender retransmission of timed-out packet
  - application-layer input = application-layer output:  $\lambda_{in} = \lambda_{out}$
  - transport-layer input includes *retransmissions*:  $\lambda'_{in} \geq \lambda_{in}$





Transport Layer 3-87

## Causes/costs of congestion: scenario 2

### idealization: perfect knowledge

- sender sends only when router buffers available




Transport Layer 3-88

## Causes/costs of congestion: scenario 2

### Idealization: known loss

packets can be lost, dropped at router due to full buffers

- sender only resends if packet *known* to be lost





Transport Layer 3-89

## Causes/costs of congestion: scenario 2

### Idealization: known loss

packets can be lost, dropped at router due to full buffers

- sender only resends if packet *known* to be lost




Transport Layer 3-90

## Causes/costs of congestion: scenario 2

### Realistic: *duplicates*


- ❖ packets can be lost, dropped at router due to full buffers
- ❖ sender times out prematurely, sending **two** copies, both of which are delivered



## Causes/costs of congestion: scenario 2

### Realistic: *duplicates*

- ❖ packets can be lost, dropped at router due to full buffers
- ❖ sender times out prematurely, sending **two** copies, both of which are delivered



### “costs” of congestion:


- ❖ more work (retrans) for given “goodput”
- ❖ unneeded retransmissions: link carries multiple copies of pkt
  - decreasing goodput

Transport Layer 3-92



## Causes/costs of congestion: scenario 3

- ❖ four senders
- ❖ multihop paths
- ❖ timeout/retransmit

**Q:** what happens as  $\lambda_{in}$  and  $\lambda_{in}'$  increase?  
**A:** as red  $\lambda_{in}'$  increases, all arriving blue pkts at upper queue are dropped, blue throughput  $\rightarrow 0$



## Causes/costs of congestion: scenario 3



### another “cost” of congestion:

- ❖ when packet dropped, any “upstream transmission capacity used for that packet was wasted!

Transport Layer 3-94

## Approaches towards congestion control

two broad approaches towards congestion control:

### end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay
- ❖ approach taken by TCP

### network-assisted congestion control:

- ❖ routers provide feedback to end systems
  - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
  - explicit rate for sender to send at

Transport Layer 3-95

## Case study: ATM ABR congestion control

### ABR: available bit rate:

- ❖ “elastic service”
- ❖ if sender’s path “underloaded”:
  - sender should use available bandwidth
- ❖ if sender’s path congested:
  - sender throttled to minimum guaranteed rate


### RM (resource management) cells:

- ❖ sent by sender, interspersed with data cells
- ❖ bits in RM cell set by switches (“network-assisted”)
  - **NI bit:** no increase in rate (mild congestion)
  - **CI bit:** congestion indication
- ❖ RM cells returned to sender by receiver, with bits intact

Transport Layer 3-96



## Case study: ATM ABR congestion control



- ❖ two-byte ER (explicit rate) field in RM cell
  - congested switch may lower ER value in cell
  - senders' send rate thus max supportable rate on path
- ❖ EFCI bit in data cells: set to 1 in congested switch
  - if data cell preceding RM cell has EFCI set, receiver sets CI bit in returned RM cell

Transport Layer 3-97


## Chapter 3 outline

- 3.1 transport-layer services
- 3.2 multiplexing and demultiplexing
- 3.3 connectionless transport: UDP
- 3.4 principles of reliable data transfer
- 3.5 connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- 3.6 principles of congestion control
- 3.7 TCP congestion control

Transport Layer 3-98


## TCP congestion control: additive increase multiplicative decrease

- ❖ **approach:** sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
  - **additive increase:** increase `cwnd` by 1 MSS every RTT until loss detected
  - **multiplicative decrease:** cut `cwnd` in half after loss



Transport Layer 3-99

## TCP Congestion Control: details



TCP sending rate:

- ❖ **roughly:** send `cwnd` bytes, wait RTT for ACKs, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

- ❖ sender limits transmission:


$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- ❖ `cwnd` is dynamic, function of perceived network congestion

Transport Layer 3-100

## TCP Slow Start

- ❖ when connection begins, increase rate exponentially until first loss event:
  - initially `cwnd` = 1 MSS
  - double `cwnd` every RTT
  - done by incrementing `cwnd` for every ACK received
- ❖ **summary:** initial rate is slow but ramps up exponentially fast



Transport Layer 3-101

## TCP: detecting, reacting to loss


- ❖ loss indicated by timeout:
  - `cwnd` set to 1 MSS;
  - window then grows exponentially (as in slow start) to threshold, then grows linearly
- ❖ loss indicated by 3 duplicate ACKs: TCP RENO
  - dup ACKs indicate network capable of delivering some segments
  - `cwnd` is cut in half window then grows linearly
- ❖ TCP Tahoe always sets `cwnd` to 1 (timeout or 3 duplicate acks)

Transport Layer 3-102

## TCP: switching from slow start to CA

Q: when should the exponential increase switch to linear?

A: when **cwnd** gets to 1/2 of its value before timeout.

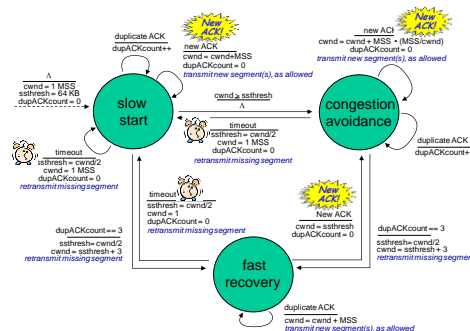


### Implementation:

- variable **ssthresh**
- on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event

Transport Layer 3-103

## Summary: TCP Congestion Control



Transport Layer 3-104

## TCP throughput

- avg. TCP thruput as function of window size, RTT?
  - ignore slow start, assume always data to send
- W: window size (measured in bytes) where loss occurs
  - avg. window size (# in-flight bytes) is  $\frac{3}{4}W$
  - avg. thruput is  $\frac{3}{4}W$  per RTT

$$\text{avg TCP thruput} = \frac{3}{4} \frac{W}{\text{RTT}} \text{ bytes/sec}$$



Transport Layer 3-105

## TCP Futures: TCP over "long, fat pipes"

- example: 1500 byte segments, 100ms RTT, want 10 Gbps throughput
- requires  $W = 83,333$  in-flight segments
- throughput in terms of segment loss probability, L [Mathis 1997]:

$$\text{TCP throughput} = \frac{1.22 \cdot \text{MSS}}{\text{RTT} \sqrt{L}}$$


→ to achieve 10 Gbps throughput, need a loss rate of  $L = 2 \cdot 10^{-10}$  — a very small loss rate!

- new versions of TCP for high-speed

Transport Layer 3-106

## TCP Fairness

**fairness goal:** if K TCP sessions share same bottleneck link of bandwidth R, each should have average rate of  $R/K$




Transport Layer 3-107

## Why is TCP fair?

two competing sessions:

- additive increase gives slope of 1, as throughput increases
- multiplicative decrease decreases throughput proportionally



Transport Layer 3-108

## Fairness (more)

### *Fairness and UDP*

- ❖ multimedia apps often do not use TCP
  - do not want rate throttled by congestion control
- ❖ instead use UDP:
  - send audio/video at constant rate, tolerate packet loss

### *Fairness, parallel TCP connections*

- ❖ application can open multiple parallel connections between two hosts
- ❖ web browsers do this
- ❖ e.g., link of rate  $R$  with 9 existing connections:
  - new app asks for 1 TCP, gets rate  $R/10$
  - new app asks for 11 TCPs, gets  $R/2$

Transport Layer 3-109

## Chapter 3: summary

- ❖ principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- ❖ instantiation, implementation in the Internet
  - UDP
  - TCP

### next:

- ❖ leaving the network “edge” (application, transport layers)
- ❖ into the network “core”

Transport Layer 3-110