


Multicast

A note on the use of these ppt slides:
The notes used in this course are substantially based on
powerpoint slides developed and copyrighted by J.F. Kurose and
K.W. Ross, 2007




*Computer Networking:
A Top Down Approach*
4th edition.
Jim Kurose, Keith Ross
Addison-Wesley, July
2007.

2-1

Broadcast Routing

- Deliver packets from source to all other nodes
- Source duplication is inefficient:



Source duplication: how does source determine recipient addresses?

2-2


In-network Duplication

- Flooding: when node receives brdcst pkt, sends copy to all neighbors
 - Problems: cycles & broadcast storm
- Controlled flooding: node only brdcsts pkt if it hasn't brdcst same packet before
 - Node keeps track of pkt ids already brdcsted
 - Or reverse path forwarding (RPF): only forward pkt if it arrived on shortest path between node and source
- Spanning tree
 - No redundant packets received by any node

2-3

Spanning Tree


- First construct a spanning tree
- Nodes forward copies only along spanning tree




2-4

Spanning Tree: Creation

- Center node
- Each node sends unicast join message to center node
 - Message forwarded until it arrives at a node already belonging to spanning tree



(a) Stepwise construction of spanning tree



(b) Constructed spanning tree


2-5

Why Multicast

- When sending same data to multiple receivers
 - Better bandwidth utilization
 - Less host/router processing
 - Quicker participation
- Application
 - Video/Audio broadcast (One sender)
 - Video conferencing (Many senders)
 - Real time news distribution
 - Interactive gaming

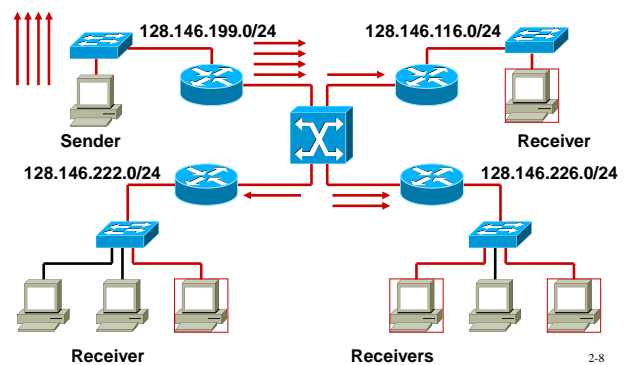
2-6

Unicast/Multicast




2-7

Unicast





2-8

Multicast



One to Many Communication

- Application-level one to many communication
- IP multicast
- Multiple unicasts




2-10

Two Major Issues

- Who are the multicast members?
- How to send the packets to the members?

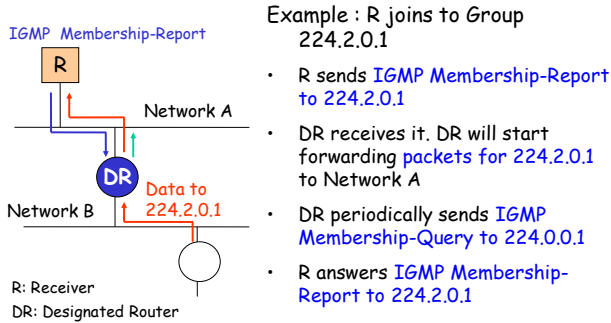
IGMP



2-11


2-12

IGMP – Joining a Group



2-13


IGMP – Leaving a Group



2-14

Multicast Routing: Problem Statement

- **Goal:** find a tree (or trees) connecting routers having local mcast group members
- **tree:** not all paths between routers used
 - **source-based:** different tree from each sender to rcvrs
 - **shared-tree:** same tree used by all group members



2-15

Approaches for Building Mcast Trees

Approaches:


- **Source-based tree:** one tree per source
 - Shortest path trees
 - Reverse path forwarding
- **Group-shared tree:** group uses one tree
 - Minimal spanning (Steiner)
 - Center-based trees

...We first look at basic approaches, then specific protocols adopting these approaches

2-16

Shortest Path Tree

- Mcast forwarding tree: tree of shortest path routes from source to all receivers
 - Dijkstra's algorithm



2-17


Reverse Path Forwarding

- Rely on router's knowledge of unicast shortest path from it to sender
- Each router has simple forwarding behavior:

if (mcast datagram received on incoming link on shortest path back to center)
then flood datagram onto all outgoing links
else ignore datagram

2-18

Reverse Path Forwarding: Example




- Result is a source-specific *reverse* SPT
 - May be a bad choice with asymmetric links

2-19

Reverse Path Forwarding: Pruning

- Forwarding tree contains subtrees with no mcast group members
 - No need to forward datagrams down subtree
 - "Prune" msgs sent upstream by router with no downstream group members



2-20

Shared-Tree: Steiner Tree

- ❑ **Steiner Tree:** minimum cost tree connecting all routers with attached group members
- ❑ Problem is NP-complete
- ❑ Excellent heuristics exists
- ❑ Not used in practice:
 - Computational complexity
 - Information about entire network needed
 - Monolithic: rerun whenever a router needs to join/leave

2-21


Center-based Trees

- ❑ Single delivery tree shared by all
- ❑ One router identified as "**center**" of tree
- ❑ To join:
 - Edge router sends unicast *join-msg* addressed to center router
 - *Join-msg* "processed" by intermediate routers and forwarded towards center
 - *Join-msg* either hits existing tree branch for this center, or arrives at center
 - Path taken by *join-msg* becomes new branch of tree for this router

2-22

Center-based Trees: an Example

- ❑ Suppose R6 chosen as center:



2-23

Internet Multicasting Routing: DVMRP

- ❑ **DVMRP:** distance vector multicast routing protocol, RFC1075
- ❑ **Flood and prune:** reverse path forwarding, source-based tree
 - RPF tree based on DVMRP's own routing tables constructed by communicating DVMRP routers
 - No assumptions about underlying unicast
 - Initial datagram to mcast group flooded everywhere via RPF
 - Routers not wanting group: send upstream prune msgs

2-24


DVMRP: continued...

- Soft state: DVMRP router periodically (1 min.) "forgets" branches are pruned:
 - Mcast data again flows down unpruned branch
 - Downstream router: re prune or else continue to receive data
- Routers can quickly regraft to tree
 - Following IGMP join at leaf
- Odds and ends
 - Commonly implemented in commercial routers
 - Mbone routing done using DVMRP

2-25

Tunneling

Q: How to connect "islands" of multicast routers in a "sea" of unicast routers?



- Mcast datagram encapsulated inside "normal" (non-multicast-addressed) datagram
- Normal IP datagram sent thru "tunnel" via regular IP unicast to receiving mcast router
- Receiving mcast router unencapsulates to get mcast datagram

2-26

PIM: Protocol Independent Multicast

- Not dependent on any specific underlying unicast routing algorithm (works with all)
- Two different multicast distribution scenarios :

Dense:

- Group members densely packed, in "close" proximity.
- Bandwidth more plentiful

Sparse:

- # Networks with group members small wrt # interconnected networks
- Group members "widely dispersed"
- Bandwidth not plentiful

2-27

Consequences of Sparse-Dense Dichotomy:

Dense:

- Group membership by routers *assumed* until routers explicitly prune
- *Data-driven* construction on mcast tree (e.g., RPF)
- Bandwidth and non-group-router processing *profligate*

Sparse:

- No membership until routers explicitly join
- *Receiver-driven* construction of mcast tree (e.g., center-based)
- Bandwidth and non-group-router processing *conservative*

2-28

PIM- Dense Mode


Flood-and-prune RPF, similar to DVMRP but

- ❑ Underlying unicast protocol provides RPF info for incoming datagram
- ❑ Less complicated (less efficient) downstream flood than DVMRP reduces reliance on underlying routing algorithm
- ❑ Has protocol mechanism for router to detect it is a leaf-node router

2-29

PIM - Sparse Mode


- ❑ Center-based approach
- ❑ Router sends *join* msg to rendezvous point (RP)
 - Intermediate routers update state and forward *join*
- ❑ After joining via RP, router can switch to source-specific tree
 - Increased performance: less concentration, shorter paths



2-30

PIM - Sparse Mode

- ❑ Sender(s):
 - Unicast data to RP, which distributes down RP-rooted tree
 - RP can extend mcast tree upstream to source
 - RP can send *stop* msg if no attached receivers
 - "no one is listening!"




2-31

ALM: Application Level Multicast

2-32


End-System Multicast

- ❑ IP multicast still is not widely deployed
 - Technical and business challenges
 - Should multicast be a *network-layer service*?
- ❑ Multicast tree of end hosts
 - Allow end hosts to form their own multicast tree
 - Hosts receiving the data help forward to others




2-33

Unicast Emulation of Multicast



2-34

IP Multicast



- No duplicate packets
- Highly efficient bandwidth usage

Key Architectural Decision: Add support for multicast in IP layer


2-35

Key Concerns with IP Multicast

- ❑ Scalability with number of groups
 - Routers maintain *per-group state*
 - Analogous to per-flow state for QoS guarantees
 - Aggregation of multicast addresses is complicated
- ❑ Supporting higher level functionality is difficult
 - IP Multicast: *best-effort multi-point delivery* service
 - End systems responsible for handling higher level functionality
 - Reliability and congestion control for IP Multicast complicated
- ❑ Deployment is difficult and slow
 - ISP's reluctant to turn on IP Multicast

2-36

End System Multicast



Potential Benefits

- Scalability
 - Routers do not maintain per-group state
 - End systems do, but they participate in very few groups
- Easier to deploy
- Potentially simplifies support for higher level functionality
 - Leverage computation and storage of end systems
 - For example, for buffering packets, transcoding, ACK aggregation
 - Leverage solutions for unicast congestion control and reliability




Design Questions

- Is End System Multicast Feasible?
- Target applications with **small and sparse groups**
- How to Build Efficient Application-Layer Multicast "Tree" or Overlay Network?
 - **Narada**: A distributed protocol for constructing efficient overlay trees among end systems
 - Simulation and Internet evaluation results to demonstrate that Narada can achieve good performance


2-39

Performance Concerns



What is an Efficient Overlay Tree?

- The delay between the source and receivers is small
- Ideally,
 - The number of redundant packets on any physical link is low
- Heuristic used:
 - Every member in the tree has a small degree
 - Degree chosen to reflect bandwidth of connection to Internet




Why is Self-Organization Hard?

- Dynamic changes in group membership
 - Members may join and leave dynamically
 - Members may die
- Limited knowledge of network conditions
 - Members do not know delay to each other when they join
 - Members probe each other to learn network related information
 - Overlay must **self-improve** as more information available
- Dynamic changes in network conditions
 - Delay between members may vary over time due to congestion

2-42

Narada Design

- | | |
|---------------|---|
| Step 1 | <p>"Mesh": Richer overlay that may have cycles and includes all group members</p> <ul style="list-style-type: none"> • Members have low degrees • Shortest path delay between any pair of members along mesh is small |
| Step 2 | <p>Source rooted shortest delay spanning trees of mesh</p> <ul style="list-style-type: none"> • Constructed using well known routing algorithms <ul style="list-style-type: none"> - Members have low degrees - Small delay from source to receivers |




2-44

Narada Components

- Mesh Management:
 - Ensures mesh remains connected in face of membership changes
- Mesh Optimization:
 - Distributed heuristics for ensuring shortest path delay between members along the mesh is small
- Spanning tree construction:
 - Routing algorithms for constructing data-delivery trees
 - Distance vector routing, and reverse path forwarding

Optimizing Mesh Quality



- Members periodically probe other members at random
- New Link added if
 - Utility Gain of adding link > Add Threshold
- Members periodically monitor existing links
- Existing Link dropped if
 - Cost of dropping link < Drop Threshold

2-45


The Terms Defined

- Utility gain of adding a link based on
 - The number of members to which routing delay improves
 - How significant the improvement in delay to each member is
- Cost of dropping a link based on
 - The number of members to which routing delay increases, for either neighbor
- Add/Drop Thresholds are functions of:
 - Member's estimation of group size
 - Current and maximum degree of member in the mesh


2-46

Desirable Properties of Heuristics

- **Stability:** A dropped link will not be immediately re-added
- **Partition Avoidance:** A partition of the mesh is unlikely to be caused as a result of any single link being dropped




Delay improves to Stan1, CMU but marginally.
Do not add link!




Delay improves to CMU, Gatech1 and significantly.
Add link!

2-47



Used by Berk1 to reach only Gatech2 and vice versa.
Drop!!




An improved mesh !!

2-48

Performance Metrics

- Delay between members using Narada
- Stress, defined as the number of identical copies of a packet that traverse a physical link



Factors Affecting Performance

- Topology Model
 - Waxman Variant
 - Mapnet: Connectivity modeled after several ISP backbones
 - ASMap: Based on inter-domain Internet connectivity
- Topology Size
 - Between 64 and 1024 routers
- Group Size
 - Between 16 and 256
- Fanout range
 - Number of neighbors each member tries to maintain in the mesh

2-50

ESM Conclusions

- Proposed in 1989, IP Multicast is not yet widely deployed
 - Per-group state, control state complexity and scaling concerns
 - Difficult to support higher layer functionality
 - Difficult to deploy, and get ISP's to turn on IP Multicast
- Is IP the right layer for supporting multicast functionality?
- For small-sized groups, an end-system overlay approach
 - is feasible
 - has a low performance penalty compared to IP Multicast
 - has the potential to simplify support for higher layer functionality
 - allows for application-specific customizations

2-51