

# LECTURE 19: SCALING MICROBLOGGING SERVICES WITH DIVERGENT TRAFFIC DEMANDS

# Outline

2

- **Background and Related Work**
- **Measuring Availability at High Load**
- **Key Design Rationales of Cuckoo**
- **Performance Evaluation**
- **Conclusion**

# Background

**Microblogging services have become a significant new form of Internet communication utility!**

## ☐ Take Twitter as an example:

1. Less than 10 years (launched in October 2006)
2. 300+ million monthly active users;  
- userbase is still growing
3. 500+ million tweets posted per day



Twitter



Google Buzz



Plurk



Sina

## Usage of Microblogging (1)

- ❑ Originally designed as an online social network (OSN) with quick updates
- ❑ Evolved to encompass a wide variety of applications, e.g.,
  1. Predict gross revenue for movie openings
  2. Produce predictions for election results
  3. Sensors for Internet service failure
  4. Real-time warning system for earthquakes



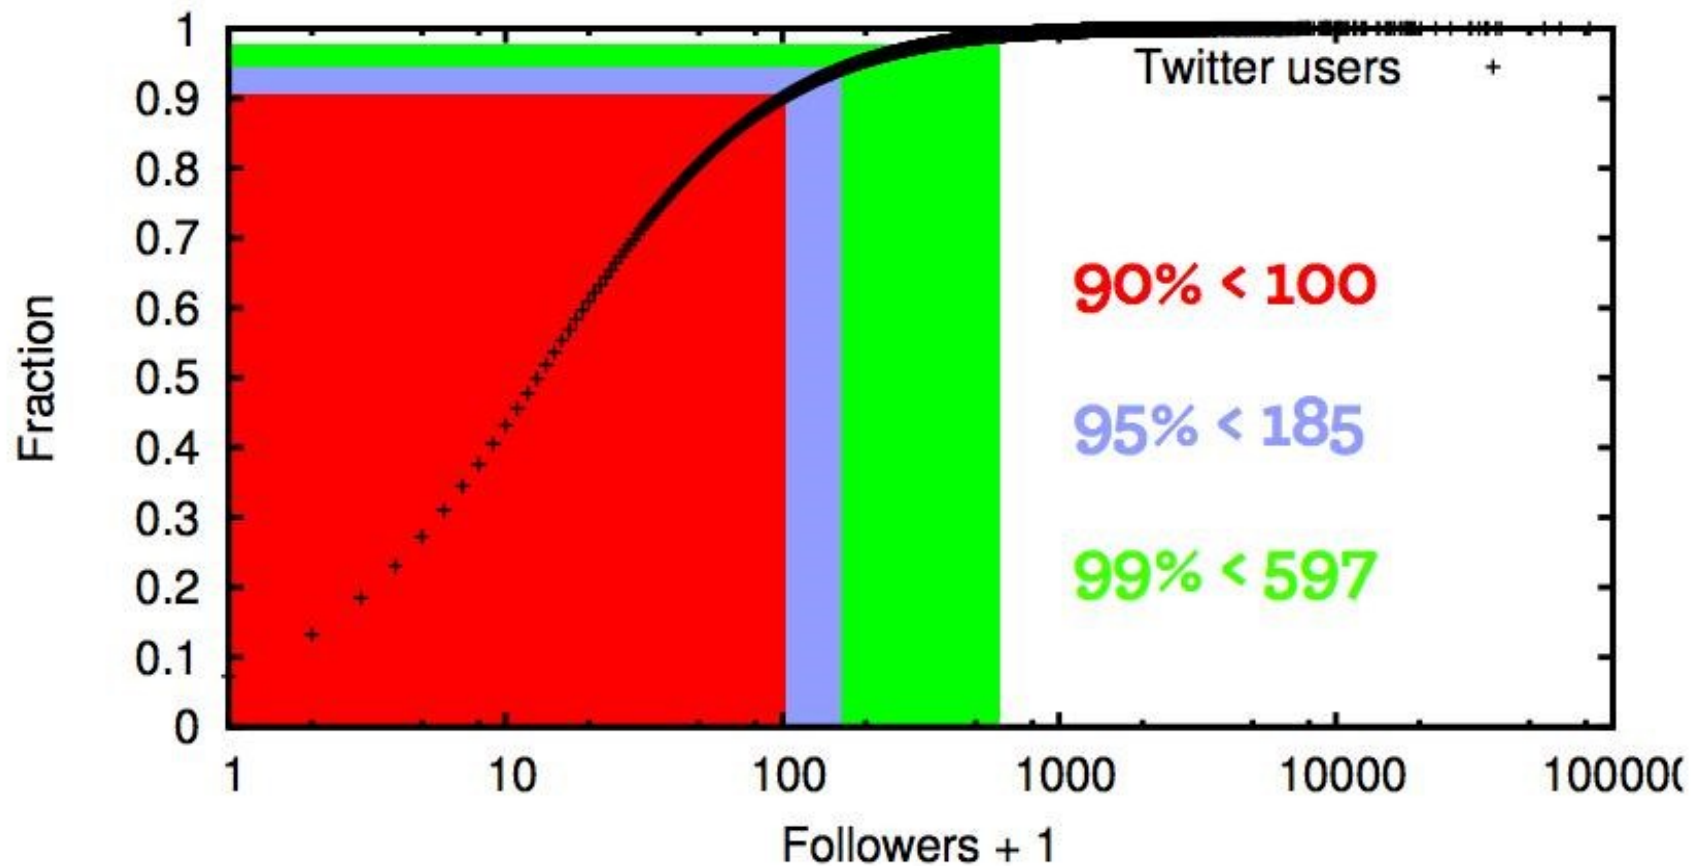
## Usage of Microblogging (2)

**Microblogging usage has also evolved significantly over time!**

- Besides as a social communication tool, much of traffic today is communication from **celebrities** and **news media** to their fans and followers.




# CDF OF TWITTER FOLLOWERS\*



\*D. R. Sandler et al., Bird of a FETHR: Open, decentralized micropublishing, IPTPS 2009.

There are a few highly-subscribed celebrities.

16.9 million subscribers












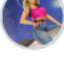


Katy Perry has 69.3 million followers

□ Twitter serves more as an **information spreading medium** than an online social network service\*.

\*H. Kwak et al., What is Twitter, a Social Network or a News Media? WWW 2010.

# Twitter Top Followers

8

Twitter users				Followers	Following	Tweets
1		KATY PERRY	@katyperry	69,391,854	156	6,481
2		Justin Bieber	@justinbieber	63,354,636	208,105	28,516
3		Barack Obama	@BarackObama	58,905,953	643,092	13,456
4		Taylor Swift	@taylorswift13	57,198,309	197	3,350
5		YouTube	@YouTube	50,643,096	889	13,607
6		Lady Gaga	@ladygaga	46,316,046	132,686	6,522
7		Justin Timberlake	@timberlake	45,052,133	106	2,938
8		Rihanna	@rihanna	44,979,193	1,176	9,688
9		Ellen DeGeneres	@TheEllenShow	43,034,454	39,008	10,522
10		Britney Spears	@britneyspears	41,762,034	398,970	4,112
11		Instagram	@instagram	38,699,862	3	6,206
12		Twitter	@twitter	38,222,422	103	1,984



## Usage of Microblogging (3)

❑ These major sources of traffic have a very tangible impact on the performance and availability of microblogging as a service.

### 1. Loss in availability

- Malicious attacks
- Hardware failures

### 2. Traffic overload and flash crowds



# Twitter's Short-Term Solutions

## 1. Per-user request and connection rate limits

- Rate limit
  - Only allows clients to make a limited number of calls in a given hour
  - Twitter: 150 requests per hour, 2,000 requests for whitelist
- Upper limit on the number of people a user could follow
  - Orkut: 1000, Flickr: 3000, Facebook: 5000,
  - Twitter: 2000 before 2009

## 2. Network usage monitoring

## 3. Doubling the capacity of internal networks

## Background

**It is clearly challenging and likely costly to scale up with demand in using the current centralized architecture.**

# Decentralized Microblogging System

## FETHR (IPTPS 2009)

- **Full decentralization**
  - Users directly contact each other via HTTP
- **Propose to use gossip for popular content propagation**
- **Cannot guarantee data delivery**
  - Some tweets cannot get to users due to user access asynchronism
- **Does not elaborate gossip component nor implementation in prototype**
- **No heterogeneous client support**
- **No user lookup service**

# Outline

13

- Background and Related Work
- **Measuring Availability at High Load**
- Key Design Rationales of Cuckoo
- Performance Evaluation
- Conclusion

# Measurement Study

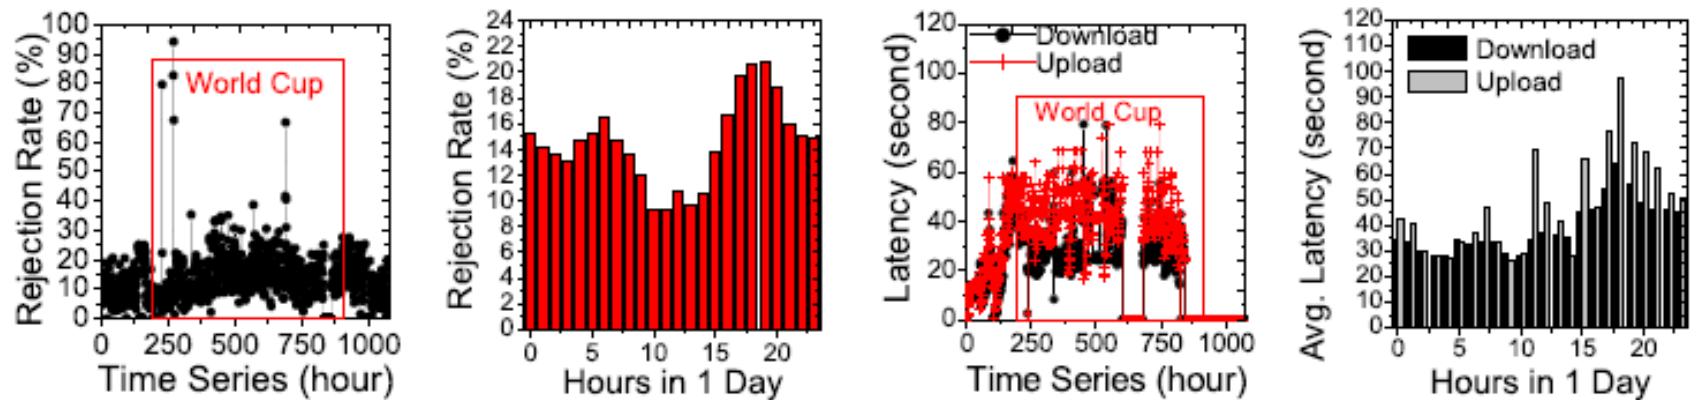
- 1. 20-day Twitter availability and performability measurement**
  - Including the period of FIFA World Cup 2010
- 2. User behavior analysis for over 3,000,000 Twitter users**
- 3. Measurement of system scalability of a generic centralized microblogging architecture**

## Availability and Performability of Twitter (1)

- **Measurement period: Jun. 4 – Jul. 18 2010**
  - Including FIFA World Cup 2010
- **Measurement place: a city in Germany**
  - The same time zone (CEST) as South Africa
- **Availability in terms of service rejection rate**
  - Randomly select a Twitter user and request for his recent 200 tweets in every 5 seconds
- **Performability\* in terms of response latency**
  - Upload latency: sending tweets to the Twitter server
  - Download latency: request the uploaded tweets from Twitter

\* P. M. Broadwell, Response Time as a Performability Metric for Online Services. Tech. Rep. UCB/CSD-04-1324, University of California at Berkeley, 2004.

## Availability and Performability of Twitter (2)



**Fig. 1.** Twitter measurement: (a), (c) Service rejection rate, response latency in time series (Jun. 4 to Jul. 18, 2010); (b), (d) Service rejection rate, response latency for 24 hours.

### □ Results

1. Twitter's availability is poor (even at normal time)
2. The flash crowd has an obvious impact on availability (service rejection rate) and performability (response latency)




# User Access Pattern Analysis (1)

- **Analysis of large-scale Twitter user trace**
  - 3,117,750 users' profile, social links, tweets
  - 4 machines with whitelisted IPs
  - Snowball crawling began with most popular 20 users reported in [1] using Twitter API
- **Consider two built-in Twitter's interaction models**
  - POST and REQUEST

[1] H. Kwak et al., What is Twitter, a Social Network or a News Media? WWW 2010.

## User Access Pattern Analysis (2)



**Fig. 2.** User access patterns: (a) # of request messages to the servers; (b) Time differences of two adjacent tweets; (c), (d) Server load in terms of received, replied messages.

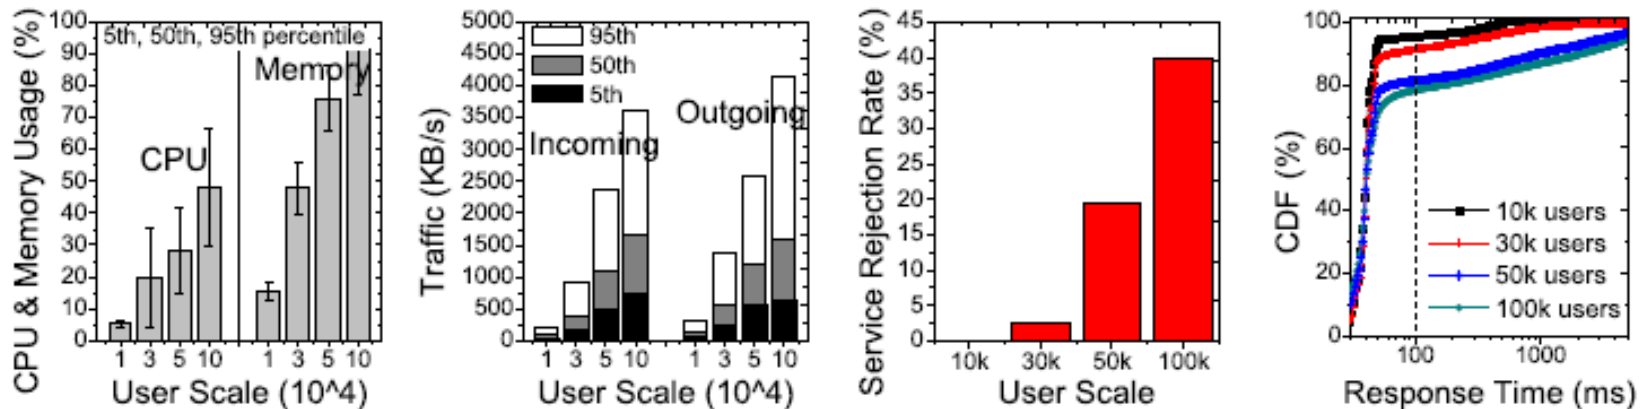
### □ Results

1. Request messages make up the dominating traffic proportion
2. There are still 50% of time differences larger than 1200 second and 20% larger than 24,000 second.
3. Although news media usage holds small proportion of incoming server load, it occupied a great proportion of outgoing server load.

# Scalability of a Generic Centralized Microblogging System (1)

- **Treat Twitter as a black box and reverse engineer its operation based on Twitter traces**
  - **Still consider POST and REQUEST as the main interaction models**
- **Each user interaction is implemented through one or more connections with centralized servers**
- **Use BFS to build four datasets for 10,000, 30,000, 50,000, 100,000 user scale and compare the server performance under different user scales**

# Scalability of a Generic Centralized Microblogging System (2)



**Fig. 3.** Scalability of the centralized microblogging architecture: (a) CPU and memory usage; (b) Traffic usage; (c) Service rejection rate; (d) Response latency.

## □ Results

1. Linear growth of CPU/memory usage and traffic usage
2. High server rejection rate when servers are overloaded
3. Long response latency when servers are overloaded

## Three Key Observations

- 1. The centralized architecture has limited scalability with the increasing number of users.**
- 2. The main server load and traffic waste are caused by polling requests.**
- 3. The social network and news media components of microblogging have distinct traffic patterns and their mix makes the system hard to scale.**

# Outline

22

- Background and Related Work
- Measuring Availability at High Load
- **Key Design Rationales of Cuckoo**
- Performance Evaluation
- Conclusion

# Decoupling the Two Components (1)

- ❑ The biggest reason that microblogging systems like Twitter do not scale:
  - They are being used as both a social network and a news media infrastructure **at the same time!**
    1. The two components have very different traffic and workload patterns with different dissemination models
      - Social network: great incoming traffic, a few followers, not very active in updating statuses
      - News media: great outgoing traffic, huge numbers of followers, highly active in posting news
    2. There is no single dissemination mechanism can really address both two at the same time.

## Decoupling the Two Components (2)




Fig. 4. Cuckoo architecture

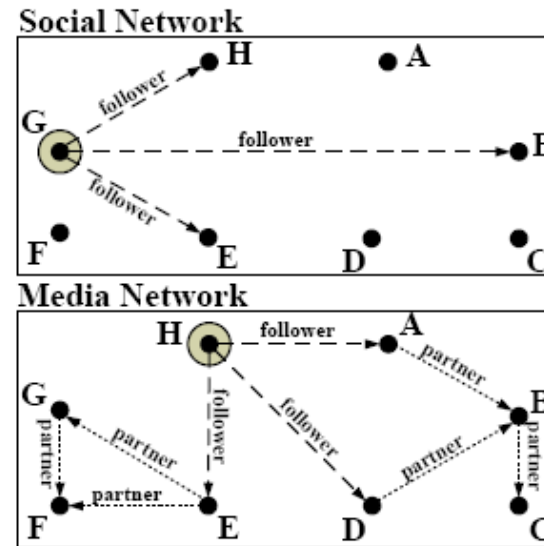


Fig. 5. Complementary content delivery

### □ Complementary delivery mechanisms

- For social network, directly push contents to followers via **unicast** (form a **social network**)
- For news media, use **gossip** to provide highly reliable and load-balanced content dissemination (**media network**)



## Combination of Server Cloud and Client Peers (1)

### ☐ **The centralized architecture is hard to scale!**

- High server load
- Performance bottleneck
- Central point of failure

### ☐ **Truly decentralized P2P systems have earned notoriety for**

- Difficult for coping with availability and consistency
- No guarantee on data delivery

## Combination of Server Cloud and Client Peers (2)

**Incorporate the advantage of both centralized and decentralized architectures:**

### ☐ **Server cloud (a small server base)**

- **Ensure high data availability**
- **Maintain asynchronous consistency**
- **Host all the user contents**

### ☐ **Cuckoo peers (client peers at network edge)**

- **Data delivery/dissemination**
  - Abandon polling -> abandon main server burden
- **Decentralized user lookup**

# Social Relations

Each user maintains three kinds of social relations:

## ❑ Followee

- User profiles

## ❑ Follower

- User profiles
- Node Handler (NH)
  - The necessary information for connection (e.g., IP, port)
- Maintain a logarithmic subset if # of followers is huge

## ❑ Partner

- Form the media network
- For each followee, a set of partner information
  - DHT-based random walk for updating



Form the **social network**

## Follow: Build Social Relations

The “Follow” operation explicitly build the followee-follower relations between user pairs.

❑ First lookup the followee’s NH and try to contact

- *Followee online*

1. Followee & Follower: build social links mutually
2. Follower: inform the server cloud of the built relation

- *Followee offline*

1. Follower: submit the willing to server cloud
2. Server cloud: check the validity and reply
3. Followee: check the inconsistency and contact the follower as compensation

# Unicast Delivery for the Social Network

- ❑ **Post a tweet => direct push to each follower**
  - Serial unicast socket
  - Via NH
- ❑ **Inform followees when changing NH**
  - Direct inform
  - Upload to the server cloud
- ❑ **Regain missing tweets in offline period**
  - From the server cloud
  - Efficient inconsistency checking
    - Based on the well-structured **timeline** (reverse chronology)

# Gossip Dissemination for the Media Network

## ❑ Enable interested users involved in the micronews dissemination process

- **Gossip-based information dissemination**
  - Scalable, resilient to network dynamics, load-balanced
  - Maintain # of partners (**fanout**) to be logarithmic of # of followers

## ❑ DHT-based partner collection

- **Announcement**
  - “Say Hello” to random destinations
  - Overhearing the hello message passing by
- **Discovery**
  - DHT-based random walk

## ❑ “Infect-and-Die” model

- **Once infected, remain infectious for one round precisely before dying**

# Support for Client Heterogeneity (1)

**Differentiate user clients into three categories:**

## ☐ **Cuckoo-Comp**

- Stable nodes
- Construct DHT and provide DHT-based user lookup
- Participate in message dissemination

## ☐ **Cuckoo-Lite**

- Lightweight clients (i.e., laptops)
- Do not join DHT
- Only participate in message dissemination

## ☐ **Cuckoo-Mobile**

- Mobile nodes
- Do not join DHT nor message dissemination

## Support for Client Heterogeneity (2)

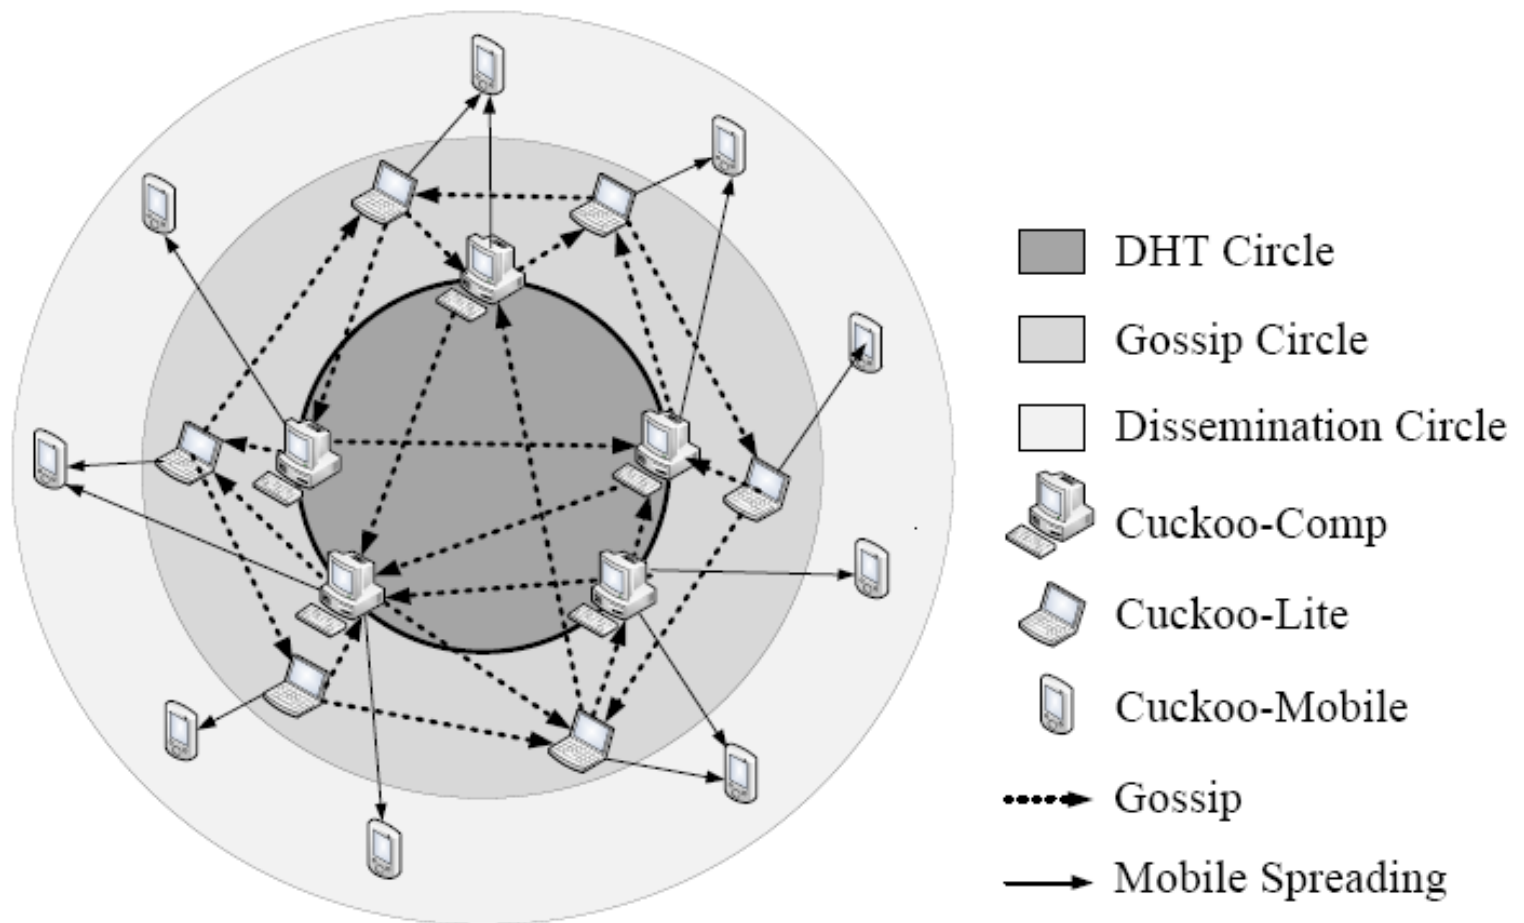


Fig. 11. Content dissemination with heterogeneous user clients



## Message Loss

**Messages may be lost due to the uncertainty of the randomness of gossip.**

**Exploit statusId to check message loss**

- ☐ **Made up of two parts: `userId` + sequence number**



- ☐ **Gaps between the sequence number of *statusIds* => message loss**

# Security Issues

- ❑ **Spam that impersonate normal users**
- ❑ **Violating message integrity by altering contents**
  - Digital signature
  - Asymmetric key cryptography
- ❑ **DoS attack and content censorship**
  - Distributed nature
  - Communication based on social/media network
  - Mark and upload later
- ❑ **Brute-force unwanted traffic**
  - Trust and reputation model
    - based on social relations


# Outline

35

- Background and Related Work
- Measuring Availability at High Load
- Key Design Rationales of Cuckoo
- **Performance Evaluation**
- Conclusion

## Dataset (1)

1. Use the raw dataset containing 3,117,750 Twitter user traces (profile, social links, tweets)
2. Filter out the 1-week part from Feb. 1 to 7, 2010
3. Use BFS to get a subset containing 30,000 users
4. Prune the irrelevant social links
5. Separate social and media with threshold 100



**Fig. 6.** Dataset: (a) CDF of users' followe{e,r} number; (b) CDF of each session duration; (c) CDF of each users' total online time.

## Dataset (2)

6. Use the OSN session duration dataset provided in [2]
7. Classify the three types of Cuckoo users according to their daily online time
  - Cuckoo-Comp: (360, 1440]
  - Cuckoo-Lite: (60, 360]
  - Cuckoo-Mobile: (0, 60]

⇒ About 50% Cuckoo peers are Cuckoo-Mobile clients.

Thus, we get the **one-week** dataset which reconstructs the part of Twitter's traffic patterns from Feb. 1 to Feb. 7, 2010.

# Implementation

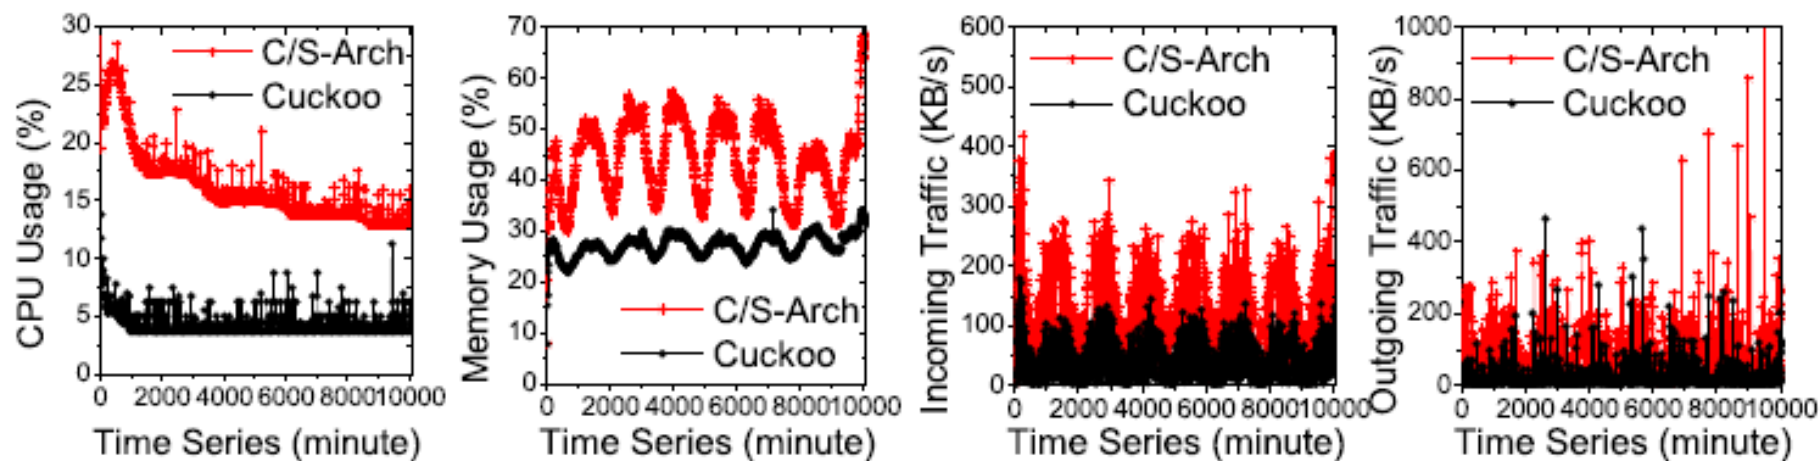
- A prototype of Cuckoo using Java comprises both the Cuckoo peer and the server cloud.
- The prototype of Cuckoo peer
  - 5000 lines of Java code
  - Three types of clients, i.e., Cuckoo-Comp, Cuckoo-Lite, and Cuckoo-Mobile
  - Use *socket* to implement end-to-end message delivery
  - Different types of application layer messages
  - Use XML for local data management
- The prototype of server cloud
  - 1500 lines of Java code
  - Use plain text files to store user information

# Deployment

- **Deploy 30,000 Cuckoo peers on 12 machines**
  - Locate these machines into 2 LANs connected by a 10 Gb/s Ethernet cable
  - Run the one-week Twitter trace
- **Deploy 4 servers to build the server cloud**
  - Let the servers share storage  
=> no inconsistency problem

# Server Cloud Performance

## - Resource Usage



**Fig. 7.** Resource usage of the server cloud in time series (from Feb. 1 to Feb. 7, 2010): (a) CPU; (b) Memory; (c) Incoming traffic; (d) Outgoing traffic.


## □ Results

1. **50% CPU usage reduction**
2. **50%/16% memory usage reduction at peak/leisure time**
3. **50% bandwidth consumed savings for both incoming and outgoing traffic**



# Server Cloud Performance

## - Message Overhead & Response Latency



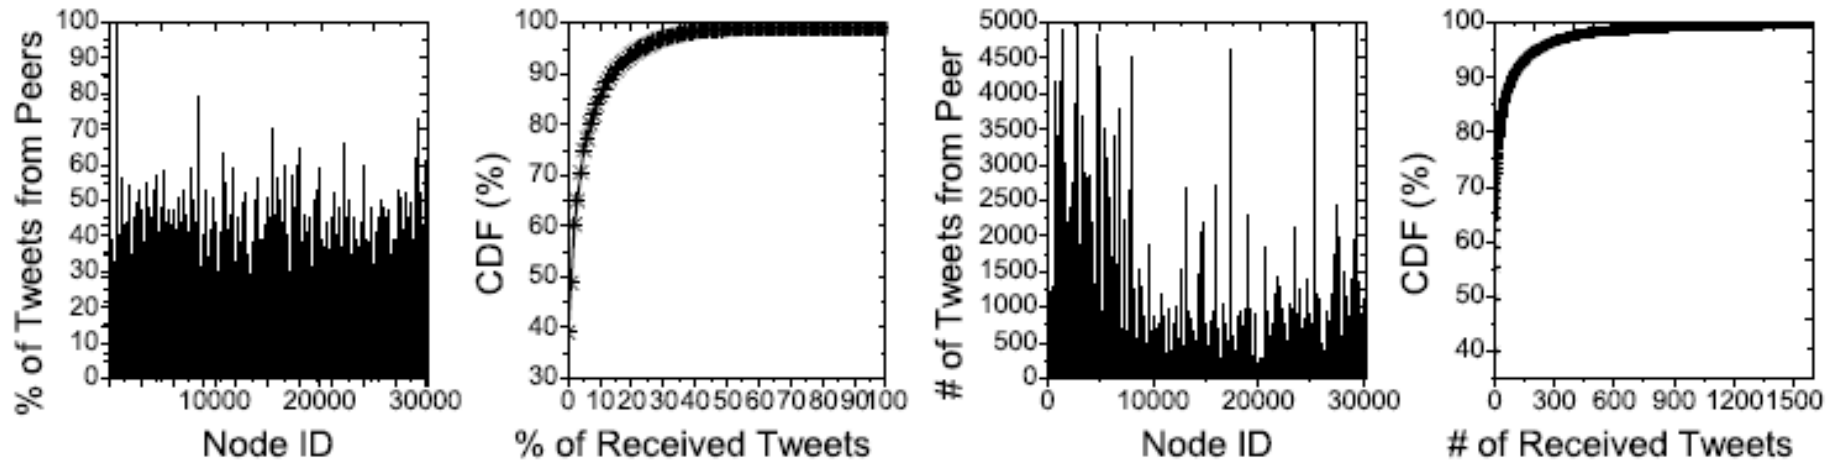
**Fig. 8.** Message overhead of the server cloud  
**Fig. 9.** Response latency of the server cloud: (a) Latency in time series; (b) CDF of the latencies.

### □ Results

1. 80% message overhead reduction
2. Smaller peak-valley difference of message overhead (140 vs 1100)
3. Requests can be satisfied within 50 ms in most time

# Cuckoo Peer Performance (1)

## - Message Sharing



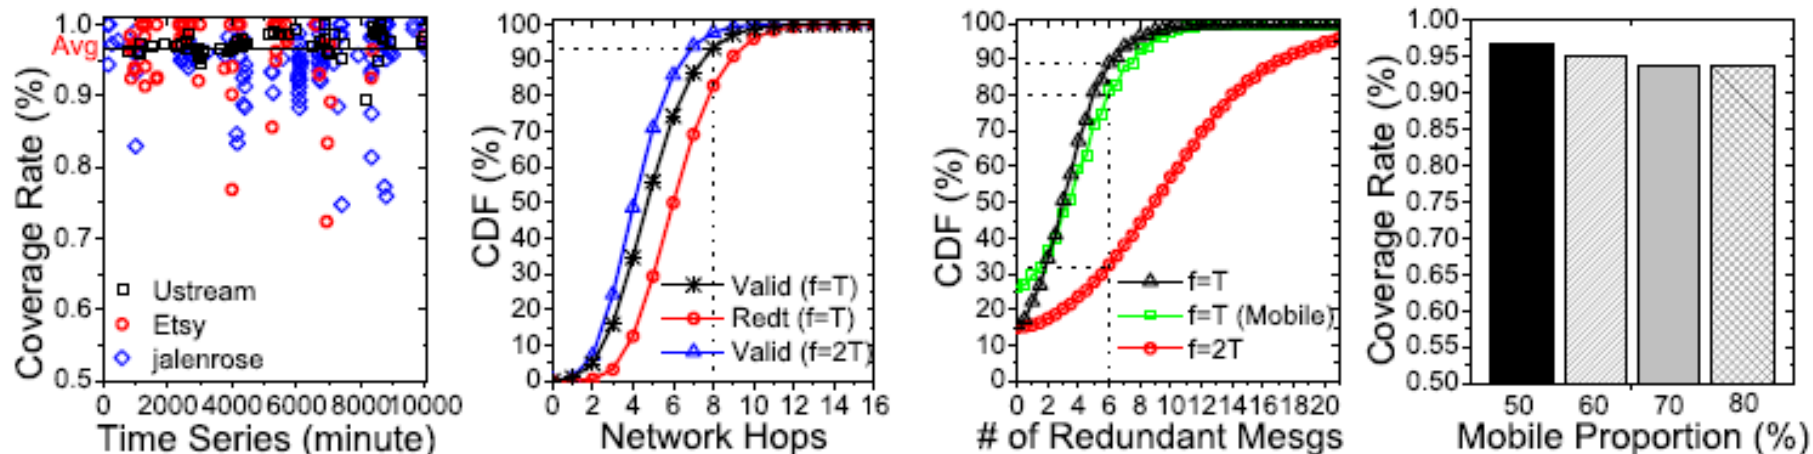
**Fig. 10.** Message sharing: (a), (b) Percentage and CDF of received tweets against overall tweets; (c), (d) Number and CDF of received tweets.

## □ Results

1. **25+% content exchanging between users**
2. **90% users receive less than 30% subscribed tweets from other peers**
  - The performance is mainly impacted by users' access behavior and users' online durations
  - The OSN dataset used leads to a pessimistic deviation

## Cuckoo Peer Performance (2)

### - Micronews Dissemination



**Fig. 11.** Micronews dissemination: (a) Coverage; (b) CDF of average network hops; (c) CDF of average redundant messages; (d) Coverage rate with different proportion of mobile nodes.

## □ Results

- a) **95+% coverage rate of content dissemination**
- b) **90% of valid micronews received are within 8 network hops**
- c) **89% of users receive less than 6 redundant tweets for one dissemination process**
- d) **Stable coverage rate with different proportion of mobile peers**

# Conclusion

- ❑ **A novel system architecture tailored for microblogging to address the scalability issues**
  - Relieve main server burden
  - Achieve scalable content delivery
  - Decoupling dual functionality components
- ❑ **A detailed measurement of Twitter**
- ❑ **A prototype implementation and trace-driven emulation over 30,000 Twitter users**
  - Notable bandwidth savings
  - Notable CPU and memory reduction
  - Good performance of content delivery/dissemination